



FASTMath Unstructured Mesh Technologies

V. Dobrev², D.A. Ibanez¹, T. Kolev², K.E. Jansen³, J.Merson⁴
O. Sahni⁴, M.S. Shephard⁴, G.M. Slota⁴, C.W. Smith⁴

¹Sandia National Laboratories

²Lawrence Livermore National Laboratory

³University of Colorado

⁴Rensselaer Polytechnic Institute



Massachusetts
Institute
of
Technology



Rensselaer



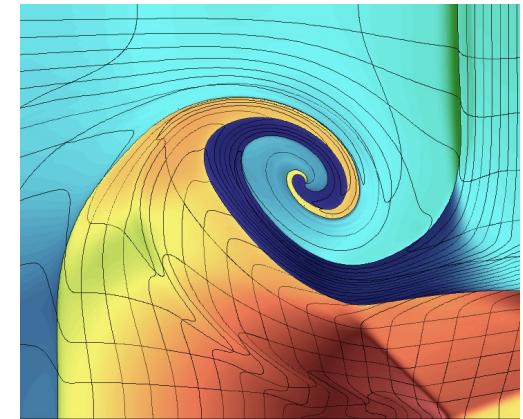
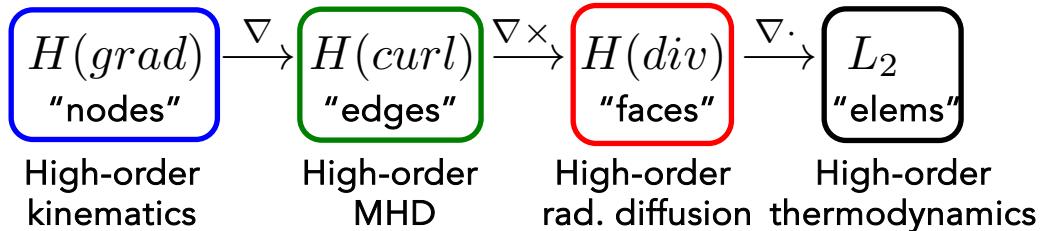
SMU



USC University of
Southern California

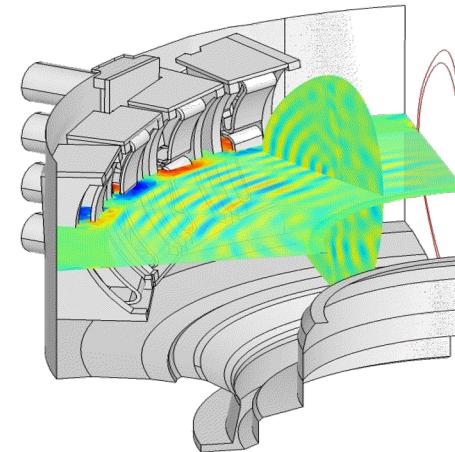
Finite elements are a good foundation for large-scale simulations on current and future architectures

- Backed by well-developed theory
- Naturally support unstructured and curvilinear grids.
- **Finite elements naturally connect different physics**



8th order Lagrangian simulation
of shock triple-point interaction

- **High-order finite elements on high-order meshes**
 - increased accuracy for smooth problems
 - sub-element modeling for problems with shocks
 - bridge unstructured/structured grids
 - bridge sparse/dense linear algebra
 - HPC utilization, FLOPs/bytes increase with the order
- **Need new (interesting!) R&D for full benefits**
 - meshing, discretizations, solvers, AMR, UQ, visualization, ...



Core-Edge tokamak EM wave
propagation

Modular Finite Element Methods (MFEM)

Flexible discretizations on unstructured grids

- Triangular, quadrilateral, tetrahedral, hexahedral, prism; volume, surface and topologically periodic meshes
- Bilinear/linear forms for: Galerkin methods, DG, HDG, DPG, IGA, ...
- Local conforming and non-conforming AMR, mesh optimization
- Hybridization and static condensation

High-order methods and scalability

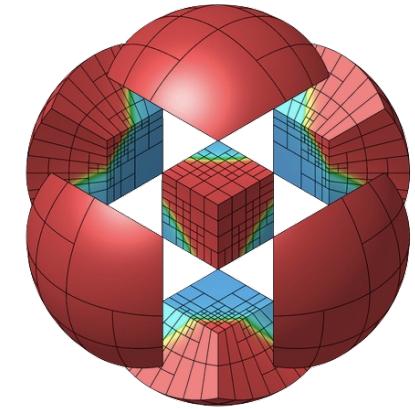
- Arbitrary-order H1, H(curl), H(div)- and L2 elements
- Arbitrary order curvilinear meshes
- MPI scalable to millions of cores + GPU accelerated
- Enables development from laptops to exascale machines.

Solvers and preconditioners

- Integrated with: HYPRE, SUNDIALS, PETSc, SLEPc, SUPERLU, VisIt, ...
- AMG solvers for full de Rham complex on CPU+GPU, geometric MG
- Time integrators: SUNDIALS, PETSc, built-in RK, SDIRK, ...

Open-source software

- Open-source (GitHub) with 114 contributors, 50 clones/day
- Part of FASTMath, ECP/CEED, xSDK, OpenHPC, E4S, ...
- 75+ example codes & miniapps: mfem.org/examples



mfem.org
(v4.5.2, Mar/2023)



Example 1 – Laplace equation

▪ Mesh

```
63 // 2. Read the mesh from the given mesh file. We can handle triangular,
64 // quadrilateral, tetrahedral, hexahedral, surface and volume meshes with
65 // the same code.
66 Mesh *mesh;
67 ifstream imesh(mesh_file);
68 if (!imesh)
69 {
70     cerr << "Can not open mesh file: " << mesh_file << '\n' << endl;
71     return 2;
72 }
73 mesh = new Mesh(imesh, 1, 1);
74 imesh.close();
75 int dim = mesh->Dimension();
76
77 // 3. Refine the mesh to increase the resolution. In this example we do
78 // 'ref_levels' of uniform refinement. We choose 'ref_levels' to be the
79 // largest number that gives a final mesh with no more than 50,000
80 // elements.
81 {
82     int ref_levels =
83         (int)floor(log(50000./mesh->GetNE()) / log(2.) / dim);
84     for (int l = 0; l < ref_levels; l++)
85         mesh->UniformRefinement();
86 }
```

▪ Finite element space

```
88 // 4. Define a finite element space on the mesh. Here we use continuous
89 // Lagrange finite elements of the specified order. If order < 1, we
90 // instead use an isoparametric/isogeometric space.
91 FiniteElementCollection *fec;
92 if (order > 0)
93     fec = new H1_FECollection(order, dim);
94 else if (mesh->GetNodes())
95     fec = mesh->GetNodes()->OwnFEC();
96 else
97     fec = new H1_FECollection(order = 1, dim);
98 FiniteElementSpace *fespace = new FiniteElementSpace(mesh, fec);
99 cout << "Number of unknowns: " << fespace->GetVSize() << endl;
```

▪ Initial guess, linear/bilinear forms

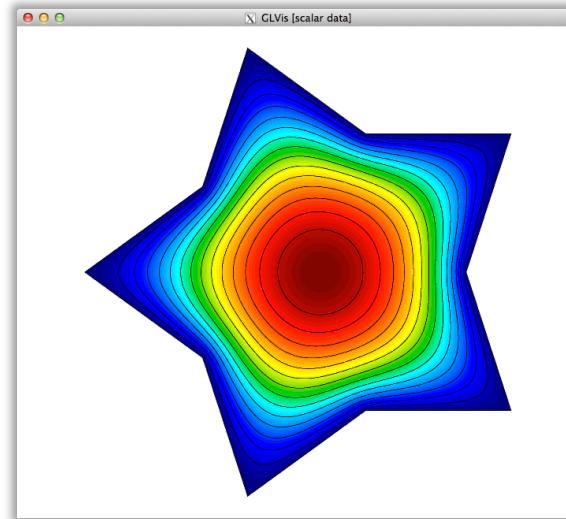
```
101 // 5. Set up the linear form b(.) which corresponds to the right-hand side of
102 // the FEM linear system, which in this case is (l,phi_i) where phi_i are
103 // the basis functions in the finite element fespace.
104 LinearForm *b = new LinearForm(fespace);
105 ConstantCoefficient one(1.0);
106 b->AddDomainIntegrator(new DomainLFIntegrator(one));
107 b->Assemble();
108
109 // 6. Define the solution vector x as a finite element grid function
110 // corresponding to fespace. Initialize x with initial guess of zero,
111 // which satisfies the boundary conditions.
112 GridFunction x(fespace);
113 x = 0.0;
114
115 // 7. Set up the bilinear form a(.,.) on the finite element space
116 // corresponding to the Laplacian operator -Delta, by adding the Diffusion
117 // domain integrator and imposing homogeneous Dirichlet boundary
118 // conditions. The boundary conditions are implemented by marking all the
119 // boundary attributes from the mesh as essential (Dirichlet). After
120 // assembly, and finalizing we extract the corresponding sparse matrix A.
121 BilinearForm *a = new BilinearForm(fespace);
122 a->AddDomainIntegrator(new DiffusionIntegrator(one));
123 a->Assemble();
124 Array<int> ess_bdr(mesh->bdr_attributes.Max());
125 ess_bdr = 1;
126 a->EliminateEssentialBC(ess_bdr, x, *b);
127 a->Finalize();
128 const SparseMatrix &A = a->SpMat();
```

▪ Linear solve

```
130 #ifndef MFEM_USE_SUITESPARSE
131 // 8. Define a simple symmetric Gauss-Seidel preconditioner and use it to
132 // solve the system Ax=b with PCG.
133 GSSmasher M(A);
134 PCG(A, M, *b, x, 1, 200, 1e-12, 0.0);
135 #else
136 // 8. If MFEM was compiled with SuiteSparse, use UMFPACK to solve the system.
137 UMFPackSolver umf_solver;
138 umf_solver.Control[UMFPACK_ORDERING] = UMFPACK_ORDERING_METIS;
139 umf_solver.SetOperator(A);
140 umf_solver.Mult(*b, x);
141#endif
```

▪ Visualization

```
152 // 10. Send the solution by socket to a GLVis server.
153 if (visualization)
154 {
155     char vishost[] = "localhost";
156     int visport = 19916;
157     socketstream sol_sock(vishost, visport);
158     sol_sock.precision(8);
159     sol_sock << "solution\n" << *mesh << x << flush;
160 }
```



- works for any mesh & any H1 order
- builds without external dependencies

Example 1 – Laplace equation

- Mesh

```
63     // 2. Read the mesh from the given mesh file. We can handle triangular,
64     //      quadrilateral, tetrahedral, hexahedral, surface and volume meshes with
65     //      the same code.
66     Mesh *mesh;
67     ifstream imesh(mesh_file);
68     if (!imesh)
69     {
70         cerr << "\nCan not open mesh file: " << mesh_file << '\n' << endl;
71         return 2;
72     }
73     mesh = new Mesh(imesh, 1, 1);
74     imesh.close();
75     int dim = mesh->Dimension();
76
77     // 3. Refine the mesh to increase the resolution. In this example we do
78     //      'ref_levels' of uniform refinement. We choose 'ref_levels' to be the
79     //      largest number that gives a final mesh with no more than 50,000
80     //      elements.
81     {
82         int ref_levels =
83             (int)floor(log(50000./mesh->GetNE()) / log(2.) / dim);
84         for (int l = 0; l < ref_levels; l++)
85             mesh->UniformRefinement();
86     }
```

Example 1 – Laplace equation

- Finite element space

```
88     // 4. Define a finite element space on the mesh. Here we use continuous
89     //      Lagrange finite elements of the specified order. If order < 1, we
90     //      instead use an isoparametric/isogeometric space.
91     FiniteElementCollection *fec;
92     if (order > 0)
93         fec = new H1_FECollection(order, dim);
94     else if (mesh->GetNodes())
95         fec = mesh->GetNodes()->OwnFEC();
96     else
97         fec = new H1_FECollection(order = 1, dim);
98     FiniteElementSpace *fespace = new FiniteElementSpace(mesh, fec);
99     cout << "Number of unknowns: " << fespace->GetVSize() << endl;
```

Example 1 – Laplace equation

- Initial guess, linear/bilinear forms

```
101 // 5. Set up the linear form b(.) which corresponds to the right-hand side of
102 // the FEM linear system, which in this case is (1,phi_i) where phi_i are
103 // the basis functions in the finite element fespace.
104 LinearForm *b = new LinearForm(fespace);
105 ConstantCoefficient one(1.0);
106 b->AddDomainIntegrator(new DomainLFIntegrator(one));
107 b->Assemble();
108
109 // 6. Define the solution vector x as a finite element grid function
110 // corresponding to fespace. Initialize x with initial guess of zero,
111 // which satisfies the boundary conditions.
112 GridFunction x(fespace);
113 x = 0.0;
114
115 // 7. Set up the bilinear form a(.,.) on the finite element space
116 // corresponding to the Laplacian operator -Delta, by adding the Diffusion
117 // domain integrator and imposing homogeneous Dirichlet boundary
118 // conditions. The boundary conditions are implemented by marking all the
119 // boundary attributes from the mesh as essential (Dirichlet). After
120 // assembly and finalizing we extract the corresponding sparse matrix A.
121 BilinearForm *a = new BilinearForm(fespace);
122 a->AddDomainIntegrator(new DiffusionIntegrator(one));
123 a->Assemble();
124 Array<int> ess_bdr(mesh->bdr_attributes.Max());
125 ess_bdr = 1;
126 a->EliminateEssentialBC(ess_bdr, x, *b);
127 a->Finalize();
128 const SparseMatrix &A = a->SpMat();
```

Example 1 – Laplace equation

- Linear solve

```
130 #ifndef MFEM_USE_SUITESPARSE
131     // 8. Define a simple symmetric Gauss-Seidel preconditioner and use it to
132     //     solve the system Ax=b with PCG.
133     GSSmooth M(A);
134     PCG(A, M, *b, x, 1, 200, 1e-12, 0.0);
135 #else
136     // 8. If MFEM was compiled with SuiteSparse, use UMFPACK to solve the system.
137     UMFPackSolver umf_solver;
138     umf_solver.Control[UMFPACK_ORDERING] = UMFPACK_ORDERING_METIS;
139     umf_solver.SetOperator(A);
140     umf_solver.Mult(*b, x);
141 #endif
```

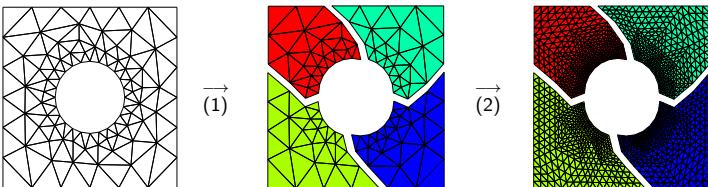
- Visualization

```
152     // 10. Send the solution by socket to a GLVis server.
153     if (visualization)
154     {
155         char vishost[] = "localhost";
156         int visport    = 19916;
157         socketstream sol_sock(vishost, visport);
158         sol_sock.precision(8);
159         sol_sock << "solution\n" << *mesh << x << flush;
160     }
```

Example 1 – parallel Laplace equation

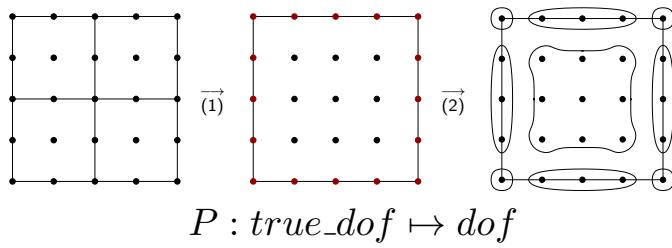
- Parallel mesh

```
101 // 5. Define a parallel mesh by a partitioning of the serial mesh. Refine
102 // this mesh further in parallel to increase the resolution. Once the
103 // parallel mesh is defined, the serial mesh can be deleted.
104 ParMesh *pmesh = new ParMesh(MPI_COMM_WORLD, *mesh);
105 delete mesh;
106 {
107     int par_ref_levels = 2;
108     for (int l = 0; l < par_ref_levels; l++)
109         pmesh->UniformRefinement();
110 }
```



- Parallel finite element space

```
122 ParFiniteElementSpace *fespace = new ParFiniteElementSpace(pmesh, fec);
```



$$P : \text{true_dof} \mapsto \text{dof}$$

- Parallel initial guess, linear/bilinear forms

```
130 ParLinearForm *b = new ParLinearForm(fespace);
131 ParGridFunction x(fespace);
132
147 ParBilinearForm *a = new ParBilinearForm(fespace);
```

- Parallel assembly

```
155 // 10. Define the parallel (hypre) matrix and vectors representing a(..),
156 // b(.) and the finite element approximation.
157 HypreParMatrix *A = a->ParallelAssemble();
158 HypreParVector *B = b->ParallelAssemble();
159 HypreParVector *X = x.ParallelAverage();
```

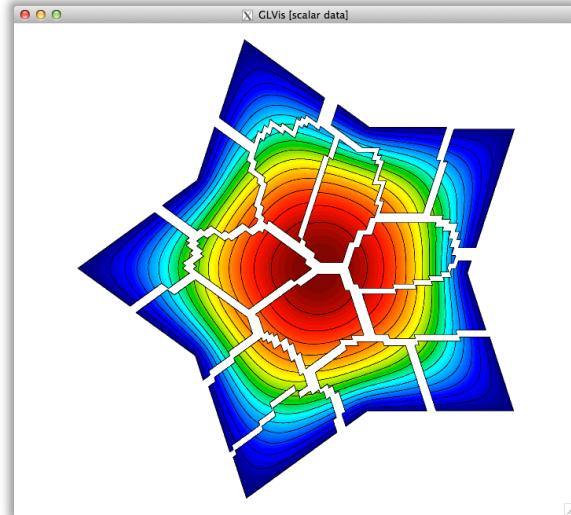
$$A = P^T a P \quad B = P^T b \quad x = P X$$

- Parallel linear solve with AMG

```
164 // 11. Define and apply a parallel PCG solver for AX=B with the BoomerAMG
165 // preconditioner from hypre.
166 HypreSolver *amg = new HypreBoomerAMG(*A);
167 HyprePCG *pcg = new HyprePCG(*A);
168 pcg->SetTol(1e-12);
169 pcg->SetMaxIter(200);
170 pcg->SetPrintLevel(2);
171 pcg->SetPreconditioner(*amg);
172 pcg->Mult(*B, *X);
```

- Visualization

```
194 // 14. Send the solution by socket to a GLVis server.
195 if (visualization)
196 {
197     char vishost[] = "localhost";
198     int visport = 19916;
199     socketstream sol_sock(vishost, visport);
200     sol_sock << "parallel" << num_procs << " " << myid << "\n";
201     sol_sock.precision(8);
202     sol_sock << "solution\n" << *pmesh << x << flush;
203 }
```



- highly scalable with minimal changes
- build depends on **hypre** and METIS

Example 1 – parallel Laplace equation

```
101 // 5. Define a parallel mesh by a partitioning of the serial mesh. Refine
102 // this mesh further in parallel to increase the resolution. Once the
103 // parallel mesh is defined, the serial mesh can be deleted.
104 ParMesh *pmesh = new ParMesh(MPI_COMM_WORLD, *mesh);
105 delete mesh;
106 {
107     int par_ref_levels = 2;
108     for (int l = 0; l < par_ref_levels; l++)
109         pmesh->UniformRefinement();
110 }

122 ParFiniteElementSpace *fespace = new ParFiniteElementSpace(pmesh, fec);
130 ParLinearForm *b = new ParLinearForm(fespace);
138 ParGridFunction x(fespace);
147 ParBilinearForm *a = new ParBilinearForm(fespace);

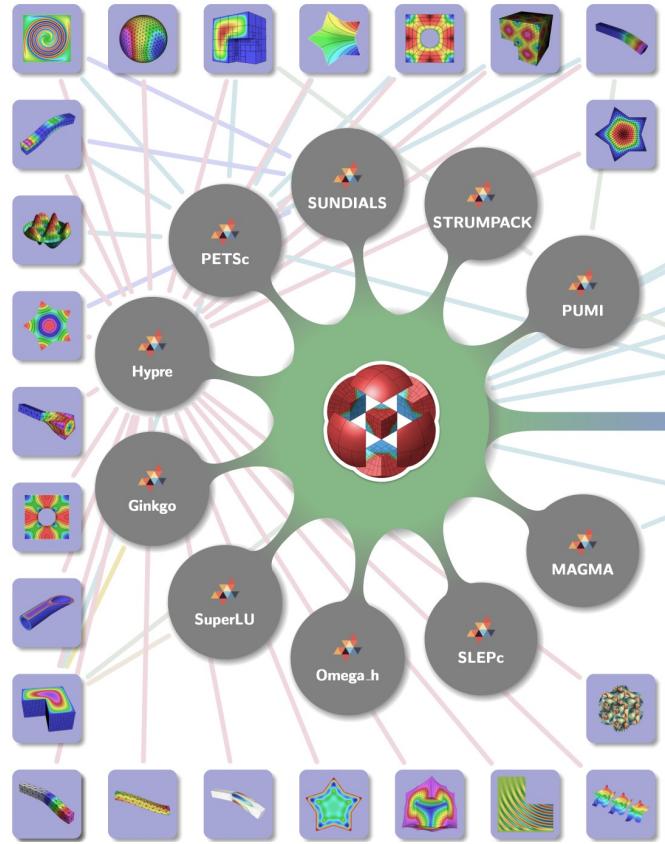
155 // 10. Define the parallel (hypre) matrix and vectors representing a(...),
156 //       b(.) and the finite element approximation.
157 HypreParMatrix *A = a->ParallelAssemble();
158 HypreParVector *B = b->ParallelAssemble();
159 HypreParVector *X = x.ParallelAverage();

164 // 11. Define and apply a parallel PCG solver for AX=B with the BoomerAMG
165 //       preconditioner from hypre.
166 HypreSolver *amg = new HypreBoomerAMG(*A);
167 HyprePCG *pcg = new HyprePCG(*A);
168 pcg->SetTol(1e-12);
169 pcg->SetMaxIter(200);
170 pcg->SetPrintLevel(2);
171 pcg->SetPreconditioner(*amg);
172 pcg->Mult(*B, *X);

200 sol_sock << "parallel " << num_procs << " " << myid << "\n";
201 sol_sock.precision(8);
202 sol_sock << "solution\n" << *pmesh << x << flush;
```

MFEM example codes: mfem.org/examples

- 40+ example codes, most with both serial + parallel versions
- Tutorials to learn MFEM features
- Starting point for new applications
- Show integration with many external packages
- Miniapps: more advanced, ready-to-use physics solvers



Demo

https://xsdk-project.github.io/MathPackagesTraining2023/lessons/mfem_convergence/

Convergence Study Source Code

The `mfem_convergence/` directory contains the `convergence.cpp` source code for solving the Poisson problem using a variety of grids and orders. You can also view the code online on [GitHub](#).

To define the system we need to solve, we need three things. First, we need to define our basis functions which live on the computational mesh.

```
// order is the FEM basis functions polynomial order
FiniteElementCollection *fec = new H1_FECollection(order, dim);

// pmesh is the parallel computational mesh
ParFiniteElementSpace *fespace = new ParFiniteElementSpace(pmesh, fec);
```

This defines a collection of H1 functions (meaning they have well-defined gradient) of a given polynomial order on a parallel computational mesh `pmesh`. Next, we need to define the integrals in Equation (5)

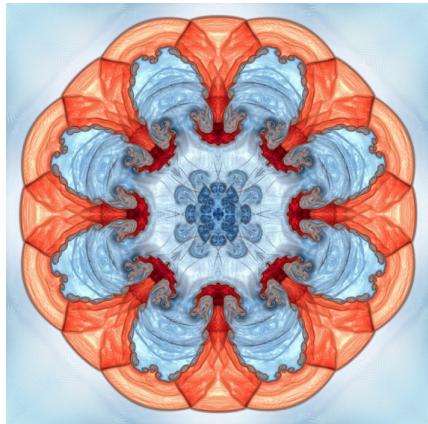
```
ParBilinearForm *a = new ParBilinearForm(fespace);
ConstantCoefficient one(1.0);
a->AddDomainIntegrator(new DiffusionIntegrator(one));
a->Assemble();
```

and Equation (6)

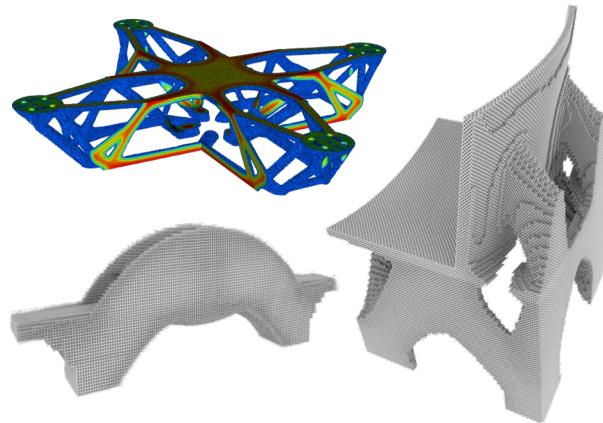
```
// f_exact is a C function defining the source
FunctionCoefficient f(f_exact);
ParLinearForm *b = new ParLinearForm(fespace);
b->AddDomainIntegrator(new DomainLFIntegrator(f));
b->Assemble();
```

This defines the matrix A and the vector b. We then solve the linear system for our solution vector x using [AMG-preconditioned](#) PCG iteration.

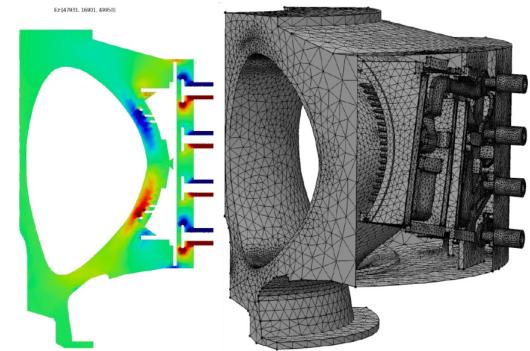
Some large-scale simulation codes powered by MFEM



Inertial confinement fusion (BLAST)



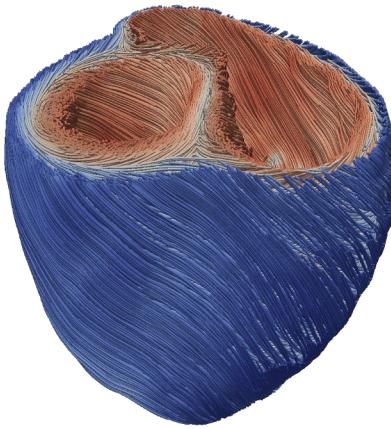
Topology optimization for additive manufacturing (LiDO)



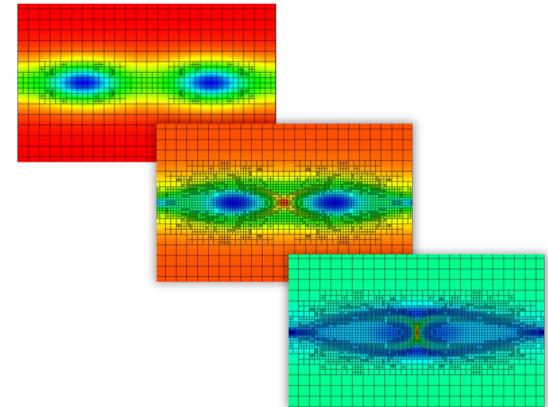
Core-edge tokamak EM wave propagation (SciDAC, RPI)



MRI modeling (Harvard Medical)



Heart modeling (Cardioid)



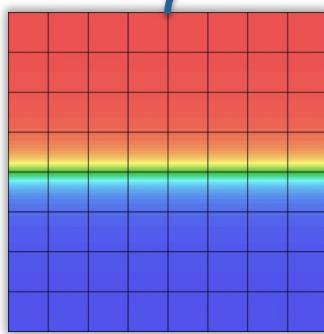
Adaptive MHD island coalescence (SciDAC, LANL)

BLAST models shock hydrodynamics using high-order FEM in both Lagrangian and Remap phases of ALE

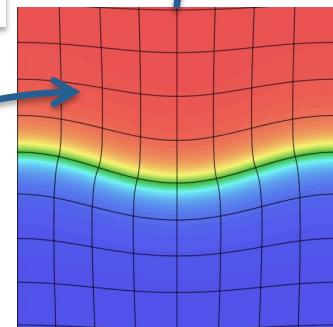
Lagrange phase

Physical time evolution

Based on physical motion



$t = 0$



$t = 1.5$

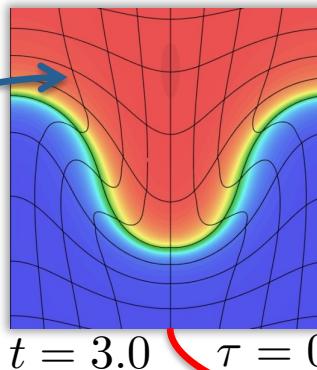
Lagrangian phase ($\vec{c} = \vec{0}$)

$$\text{Momentum Conservation: } \rho \frac{d\vec{v}}{dt} = \nabla \cdot \sigma$$

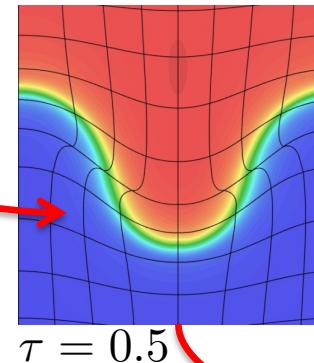
$$\text{Mass Conservation: } \frac{d\rho}{dt} = -\rho \nabla \cdot \vec{v}$$

$$\text{Energy Conservation: } \rho \frac{de}{dt} = \sigma : \nabla \vec{v}$$

$$\text{Equation of Motion: } \frac{d\vec{x}}{dt} = \vec{v}$$



$t = 3.0 \quad \tau = 0$

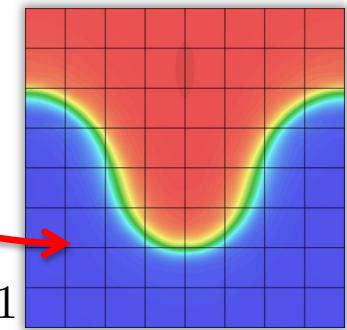


$\tau = 0.5$

Remap phase

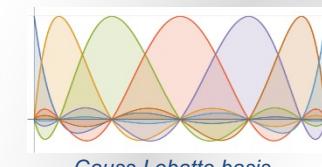
Pseudo-time evolution

Based on mesh motion



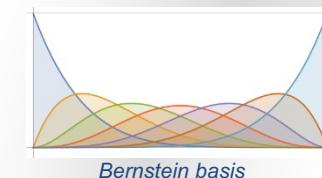
$\tau = 1$

Galerkin FEM



Gauss-Lobatto basis

Discont. Galerkin



Bernstein basis

Advection phase ($\vec{c} = -\vec{v}_m$)

$$\text{Momentum Conservation: } \frac{d(\rho \vec{v})}{d\tau} = \vec{v}_m \cdot \nabla(\rho \vec{v})$$

$$\text{Mass Conservation: } \frac{d\rho}{d\tau} = \vec{v}_m \cdot \nabla \rho$$

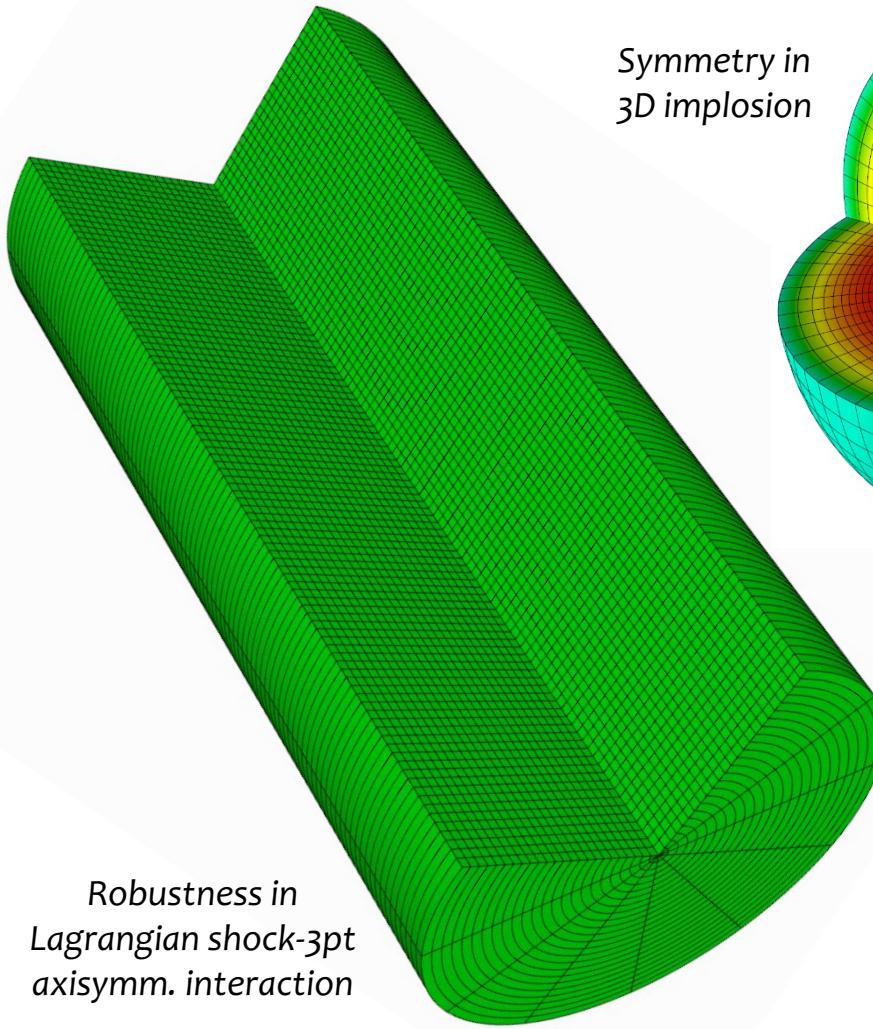
$$\text{Energy Conservation: } \frac{d(\rho e)}{d\tau} = \vec{v}_m \cdot \nabla(\rho e)$$

$$\text{Mesh velocity: } \vec{v}_m = \frac{d\vec{x}}{d\tau}$$

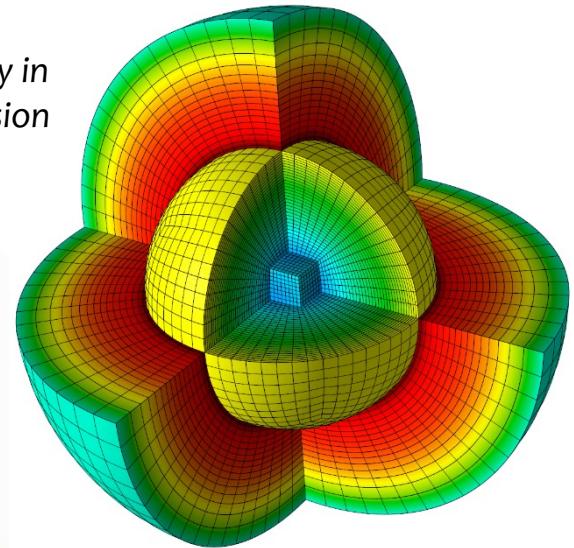
High-order finite elements lead to more accurate, robust and reliable hydrodynamic simulations



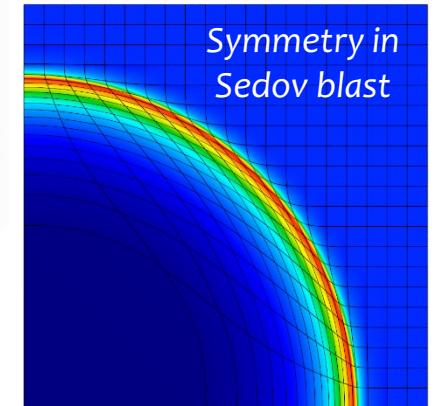
Parallel ALE for Q4 Rayleigh-Taylor instability (256 cores)



Symmetry in
3D implosion

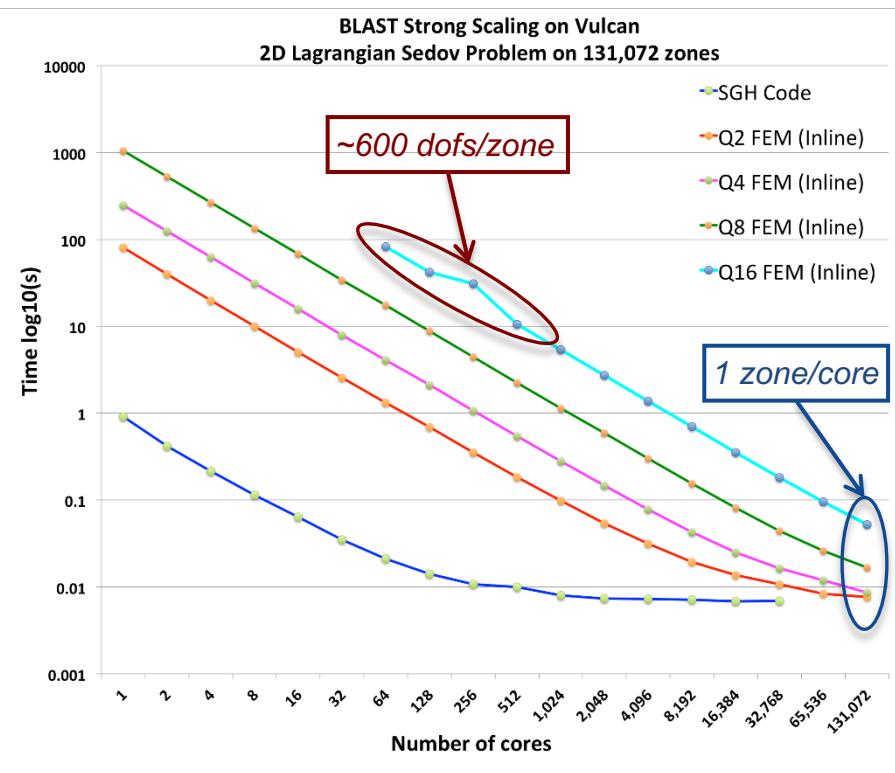


Symmetry in
Sedov blast



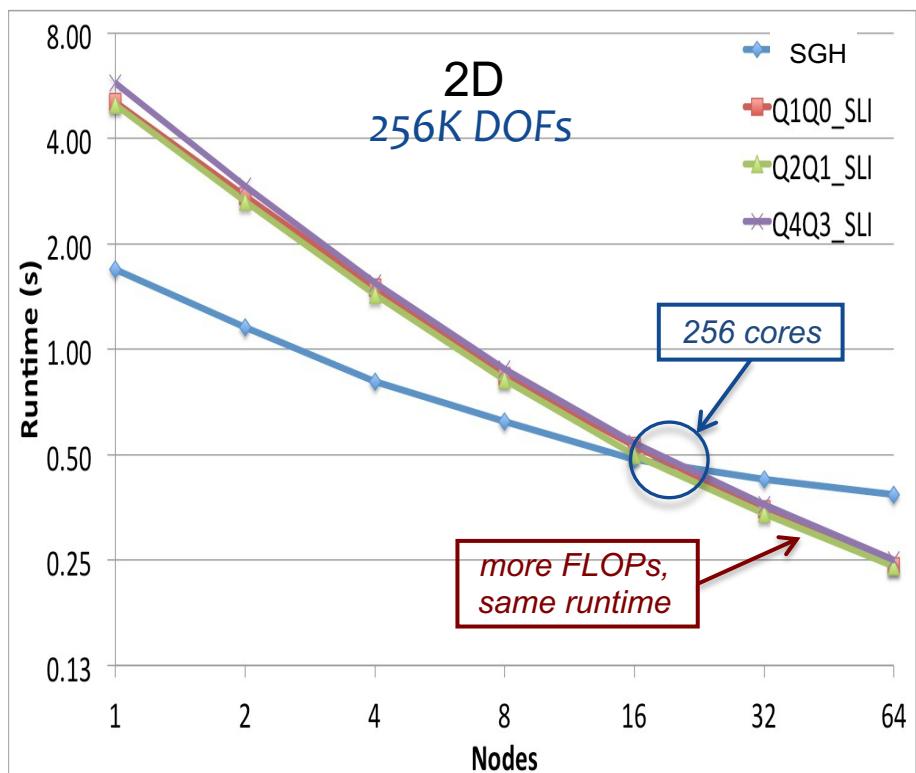
High-order finite elements have excellent strong scalability

Strong scaling, p -refinement



Finite element partial assembly

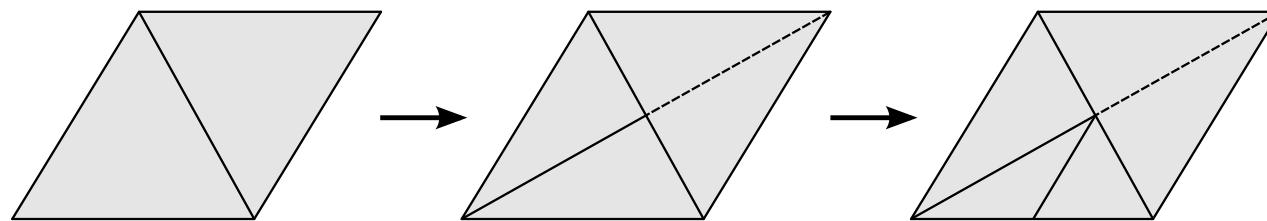
Strong scaling, fixed #dofs



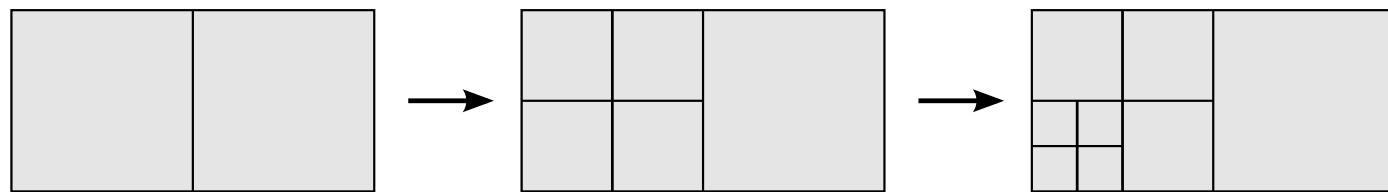
FLOPs increase faster than runtime

Conforming & Nonconforming Mesh Refinement

■ Conforming refinement



■ Nonconforming refinement



■ Natural for quadrilaterals and hexahedra

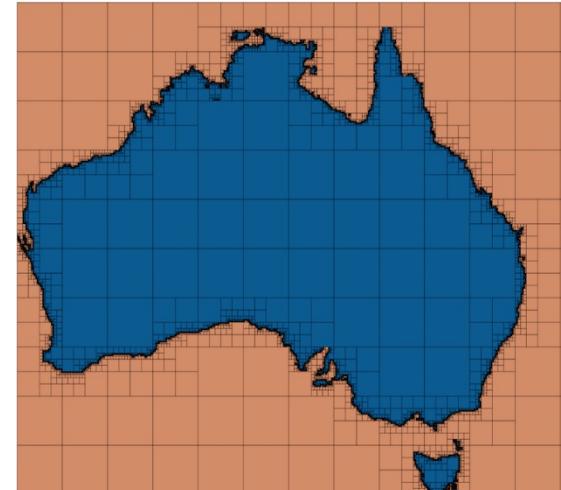
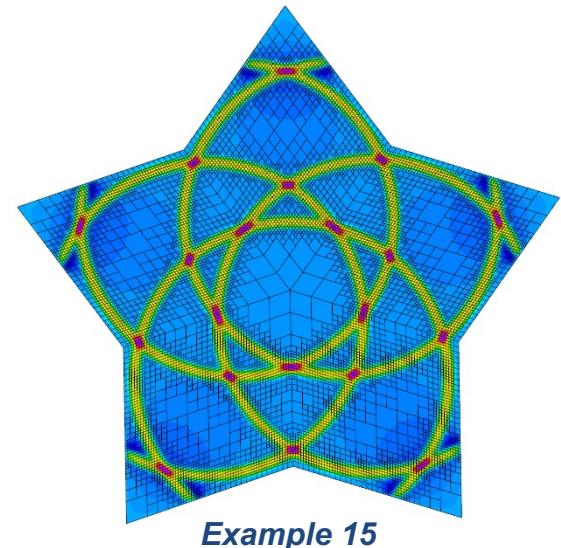
MFEM's unstructured AMR infrastructure

Adaptive mesh refinement on library level:

- Conforming local refinement on simplex meshes
- **Non-conforming refinement for quad/hex meshes**
- h-refinement with fixed p

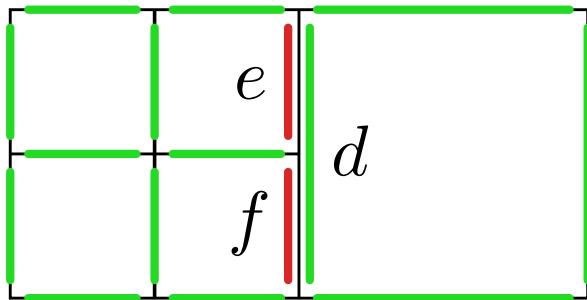
General approach:

- any high-order finite element space, H^1 , $H(\text{curl})$, $H(\text{div})$, ..., on any high-order curved mesh
- 2D and 3D
- arbitrary order hanging nodes
- anisotropic refinement
- derefinement
- serial and parallel, including parallel load balancing
- independent of the physics (easy to incorporate in applications)



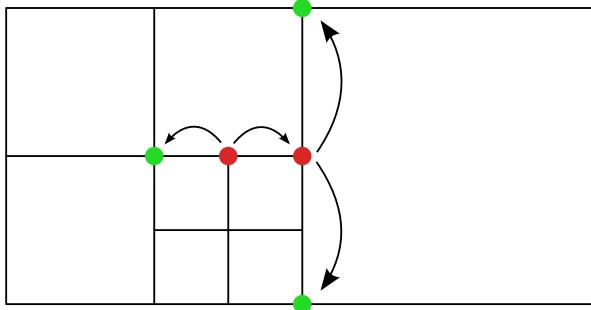
General nonconforming constraints

$H(\text{curl})$ elements



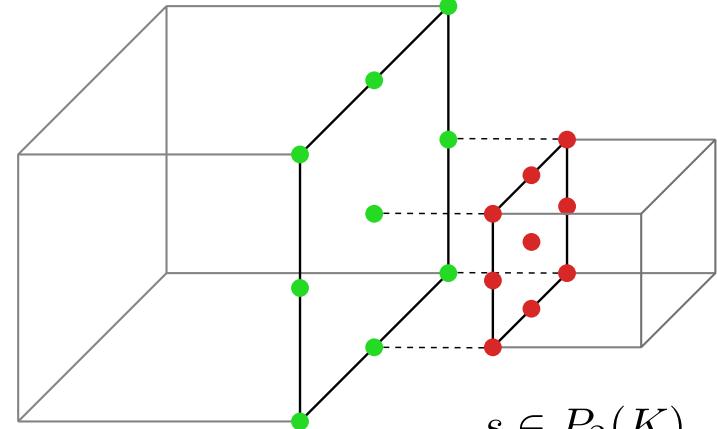
Constraint: $e = f = d/2$

Indirect constraints



More complicated in 3D...

High-order elements



$m \in P_2(K)$

Constraint: local interpolation matrix

$$s = Q \cdot m, \quad Q \in \mathbb{R}^{9 \times 9}$$

Nonconforming variational restriction

- General constraint:

$$y = Px, \quad P = \begin{bmatrix} I \\ W \end{bmatrix}.$$

x – conforming space DOFs,

y – nonconforming space DOFs (unconstrained + slave),

$$\dim(x) \leq \dim(y)$$

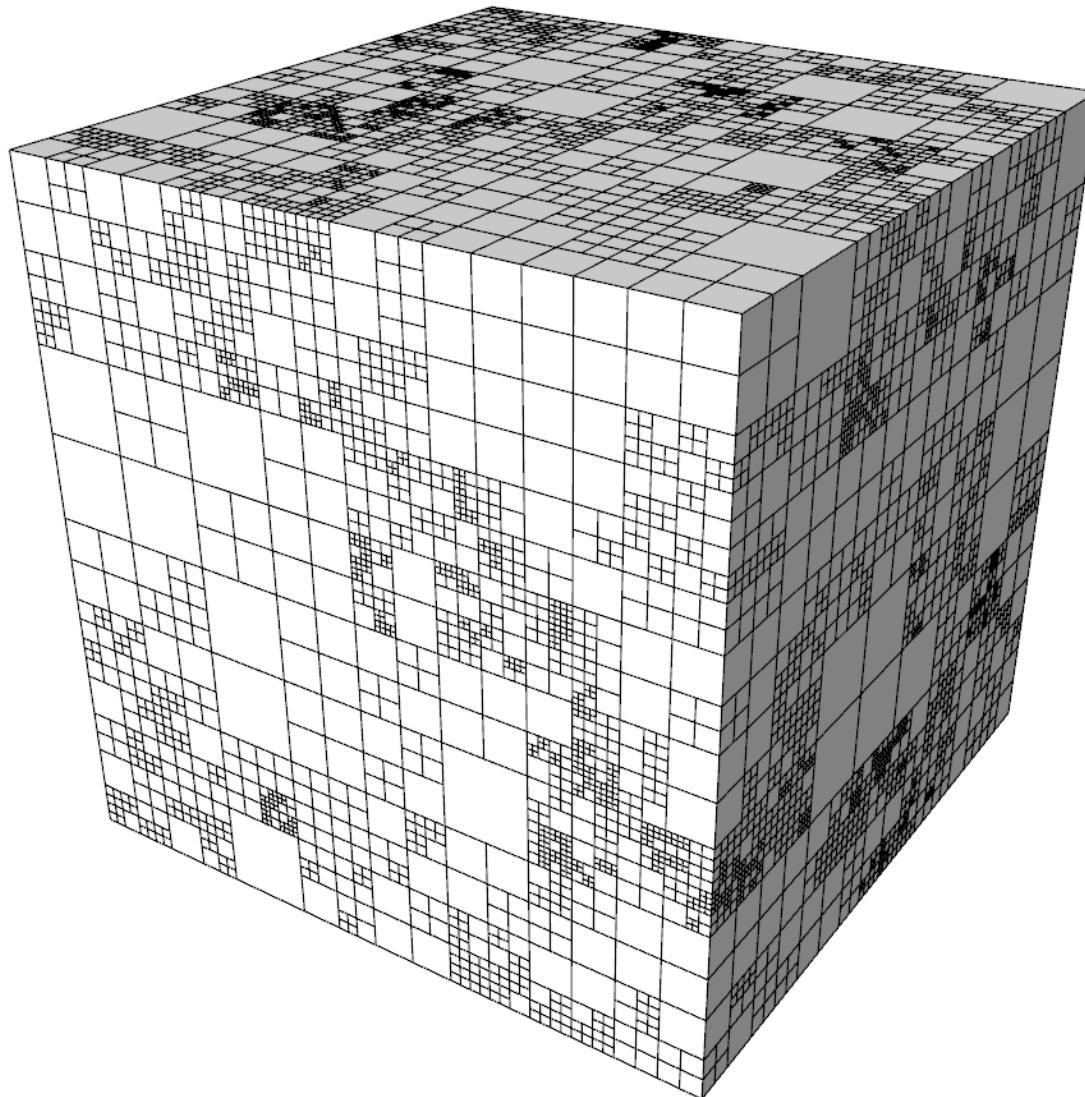
W – interpolation for slave DOFs

- Constrained problem:

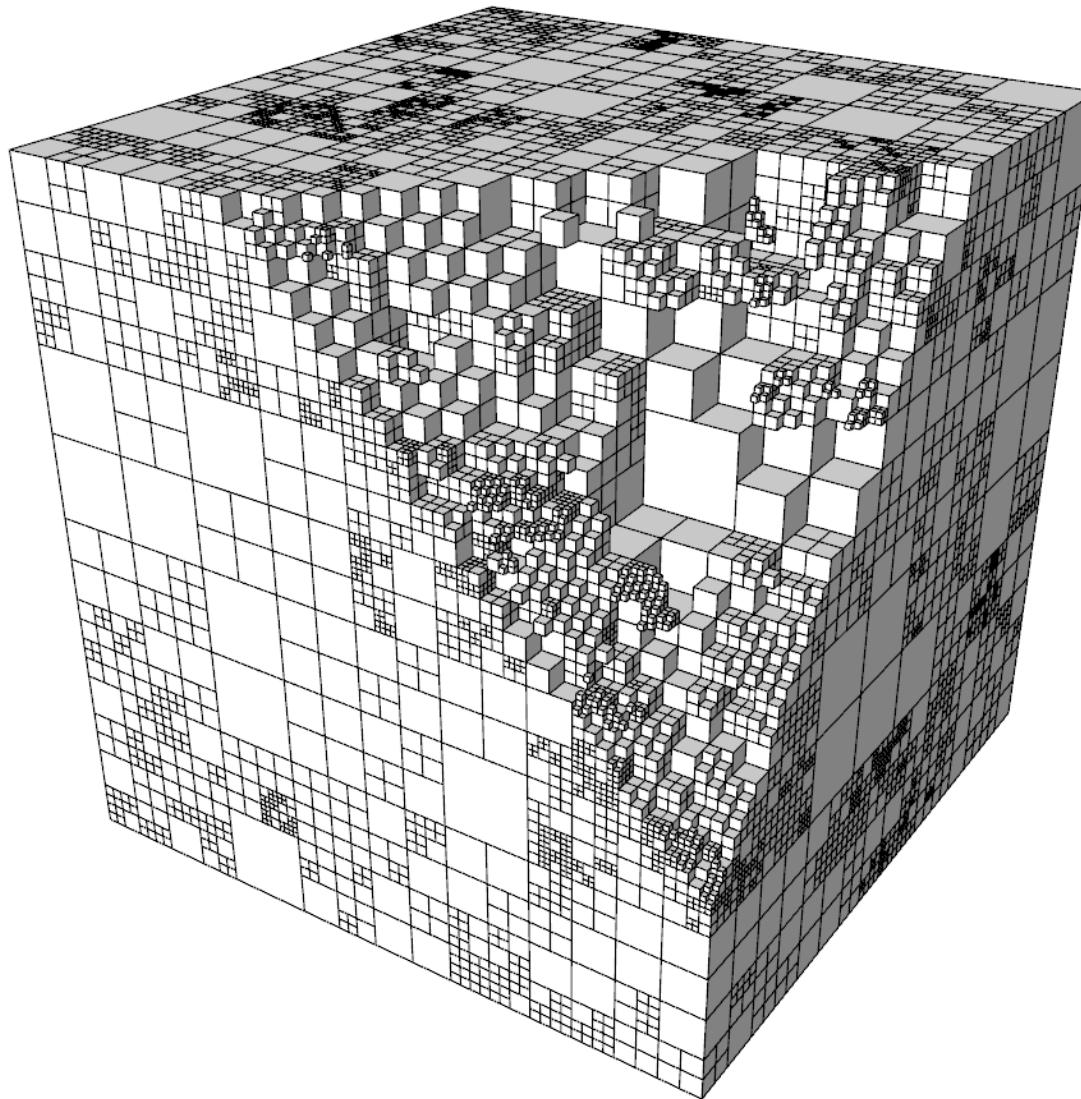
$$P^T A P x = P^T b,$$

$$y = Px.$$

Nonconforming variational restriction

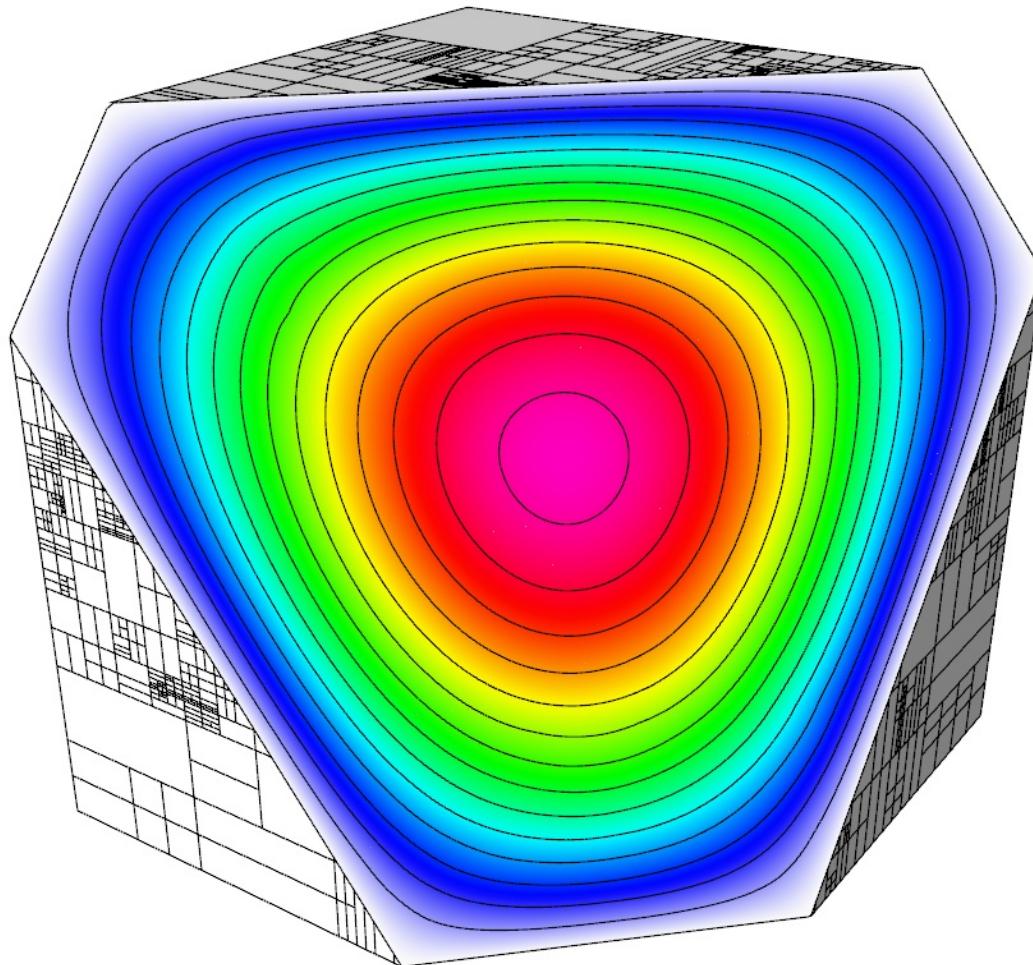


Nonconforming variational restriction



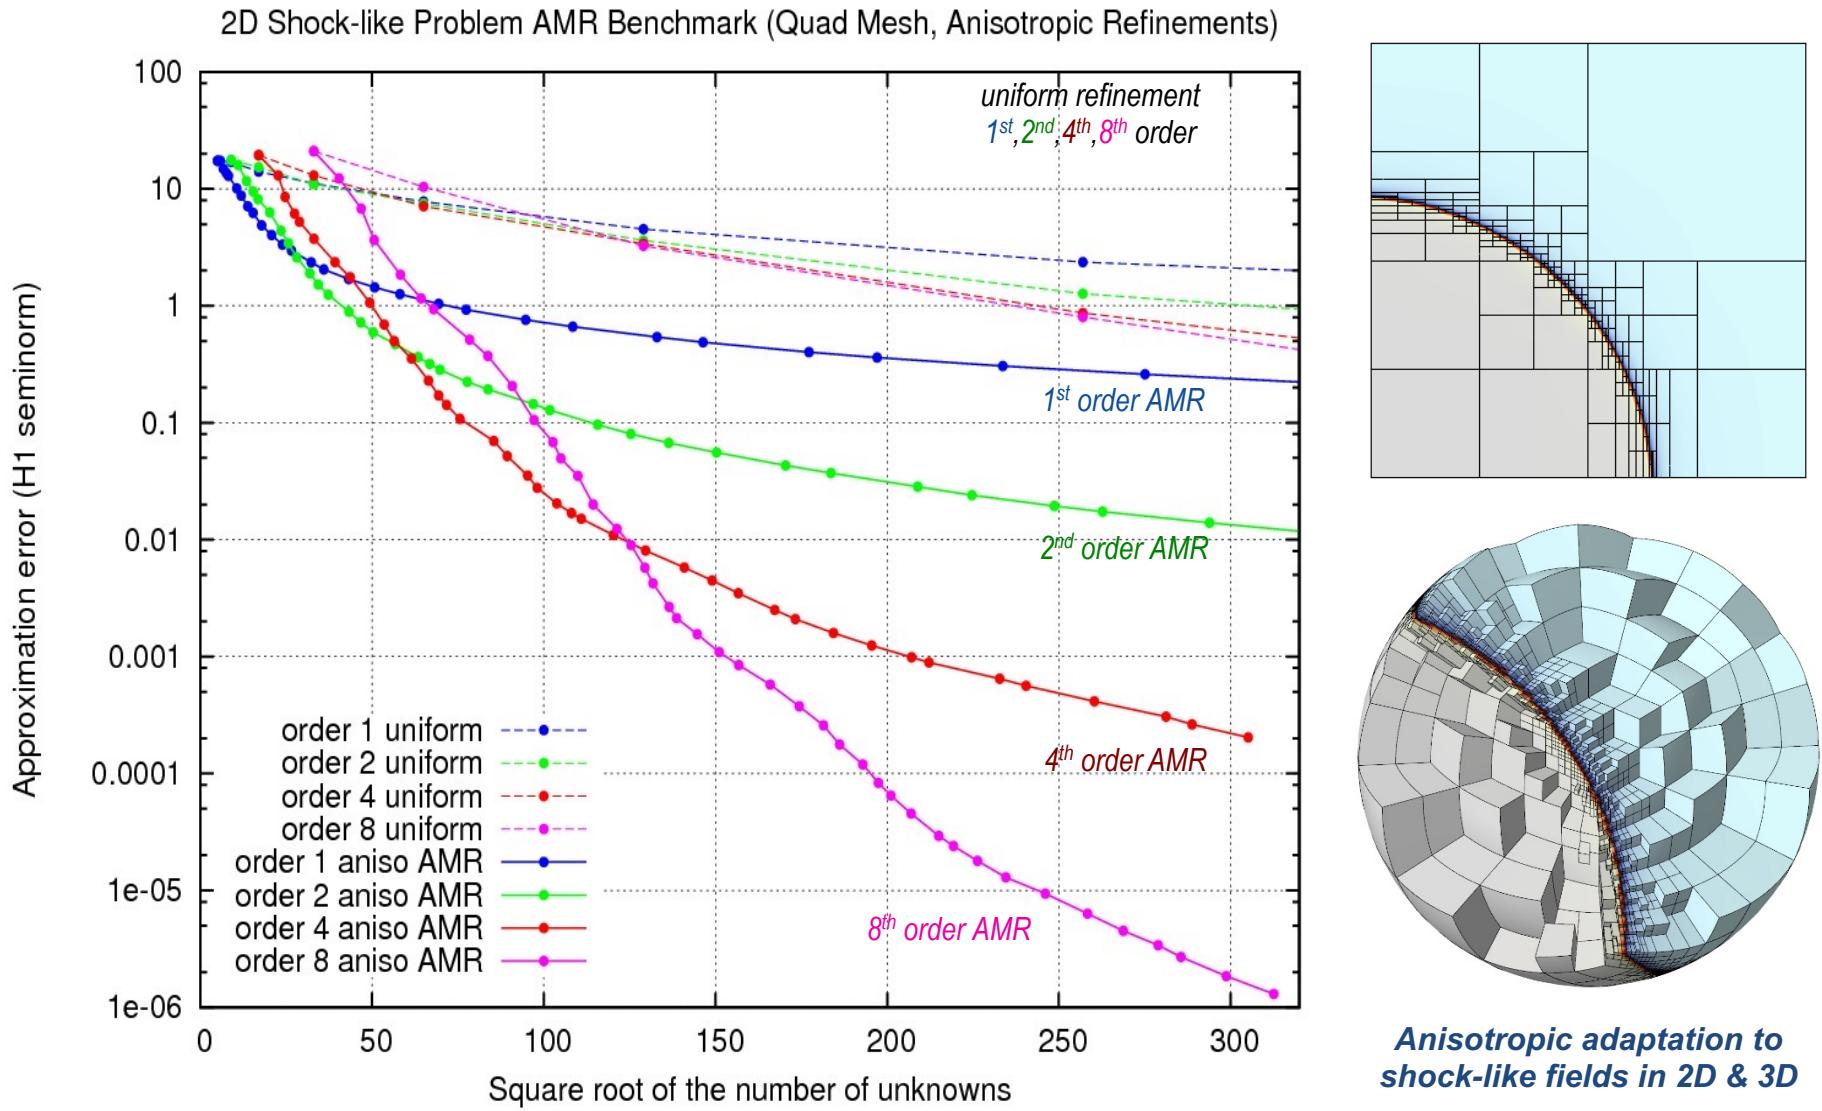
Regular assembly of A on the elements of the (cut) mesh

Nonconforming variational restriction

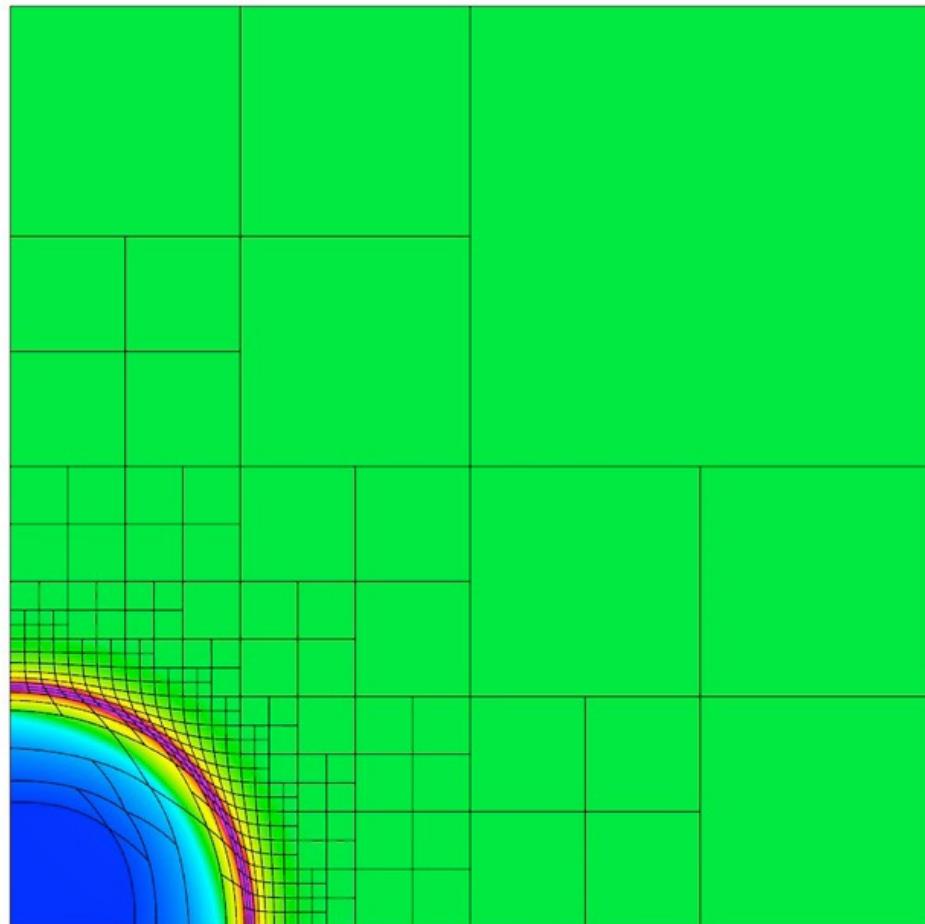


Conforming solution $y = P x$

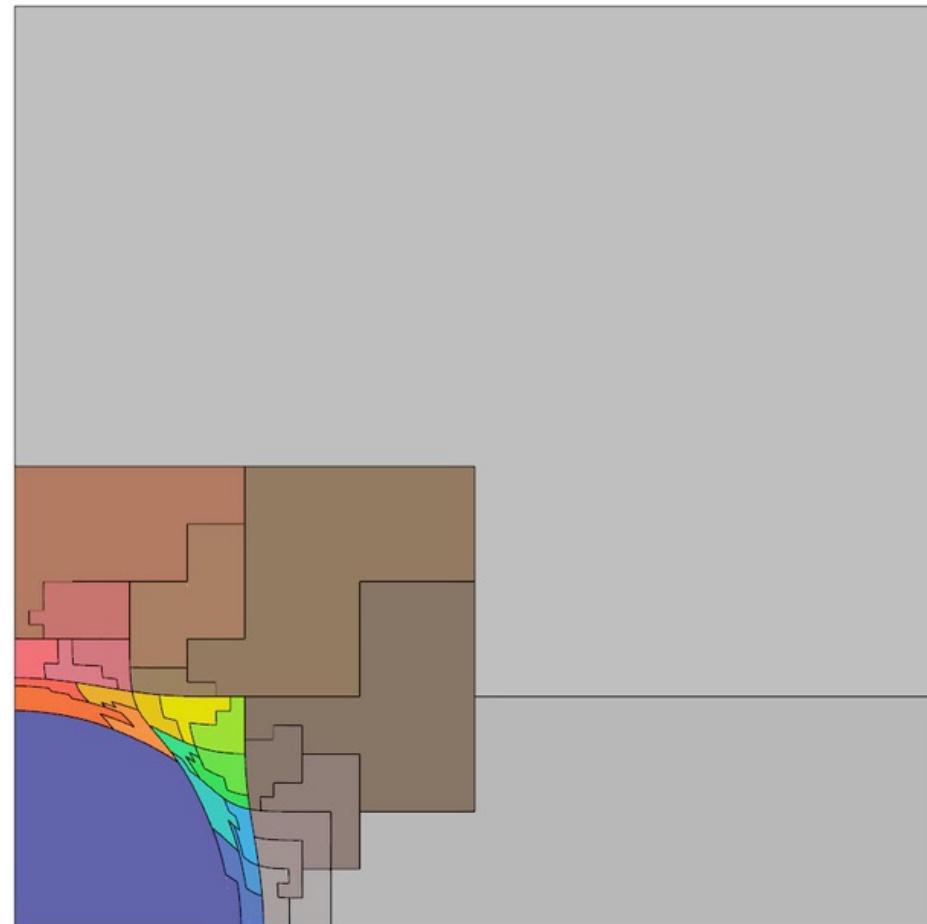
AMR = smaller error for same number of unknowns



Parallel dynamic AMR, Lagrangian Sedov problem

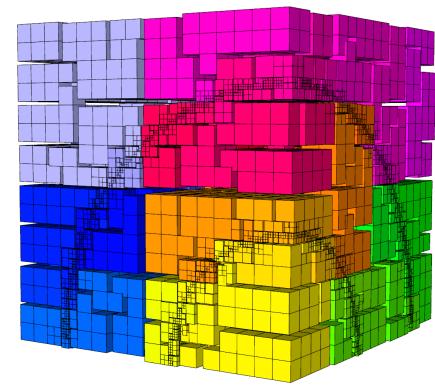
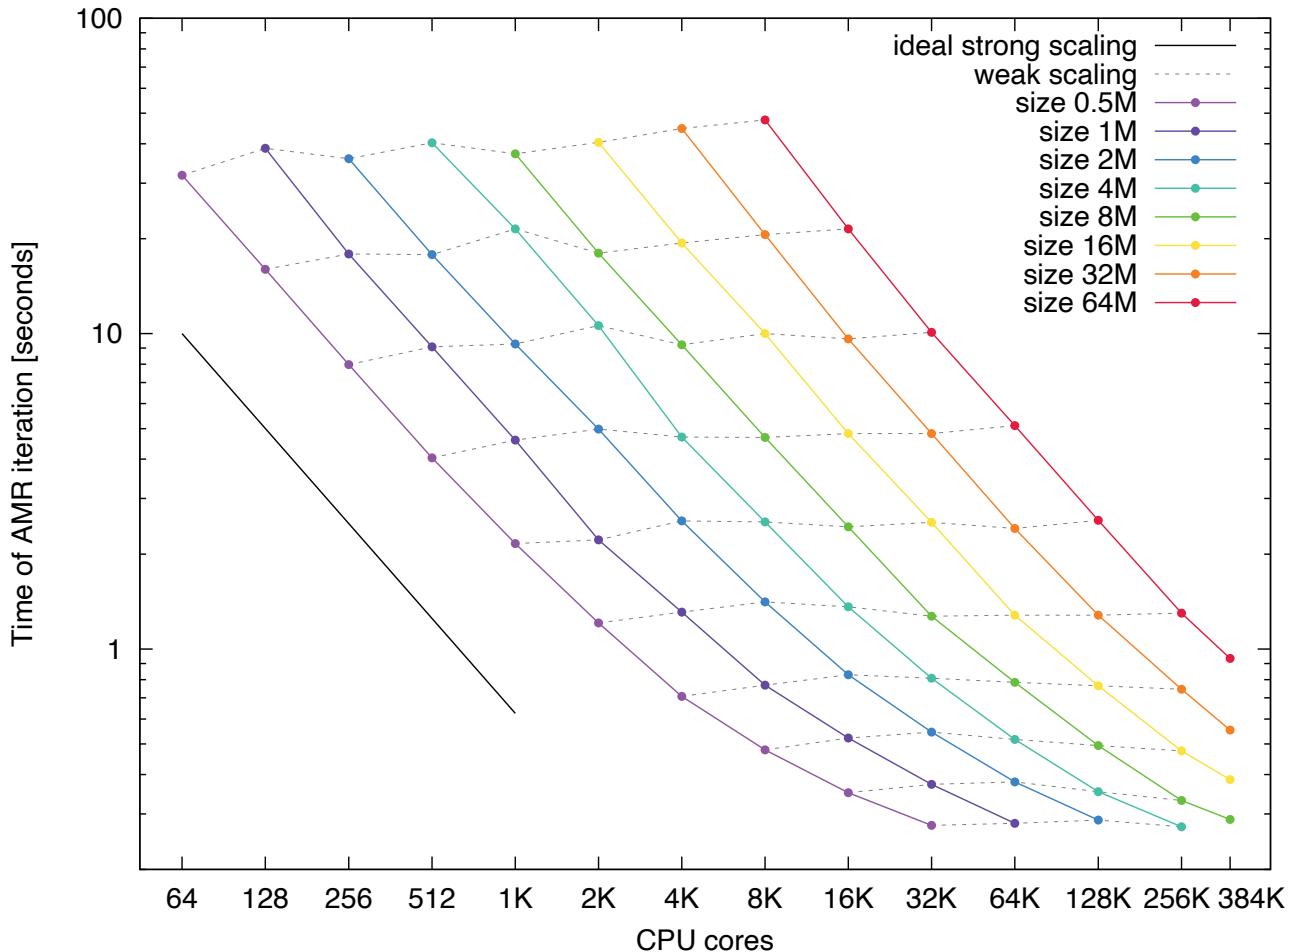


Adaptive, viscosity-based refinement and derefinement. 2nd order Lagrangian Sedov

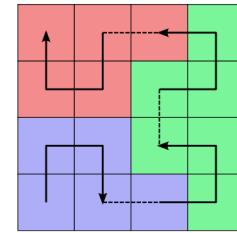


Parallel load balancing based on space-filling curve partitioning, 16 cores

Parallel AMR scaling to ~400K MPI tasks



**Parallel decomposition
(2048 domains shown)**

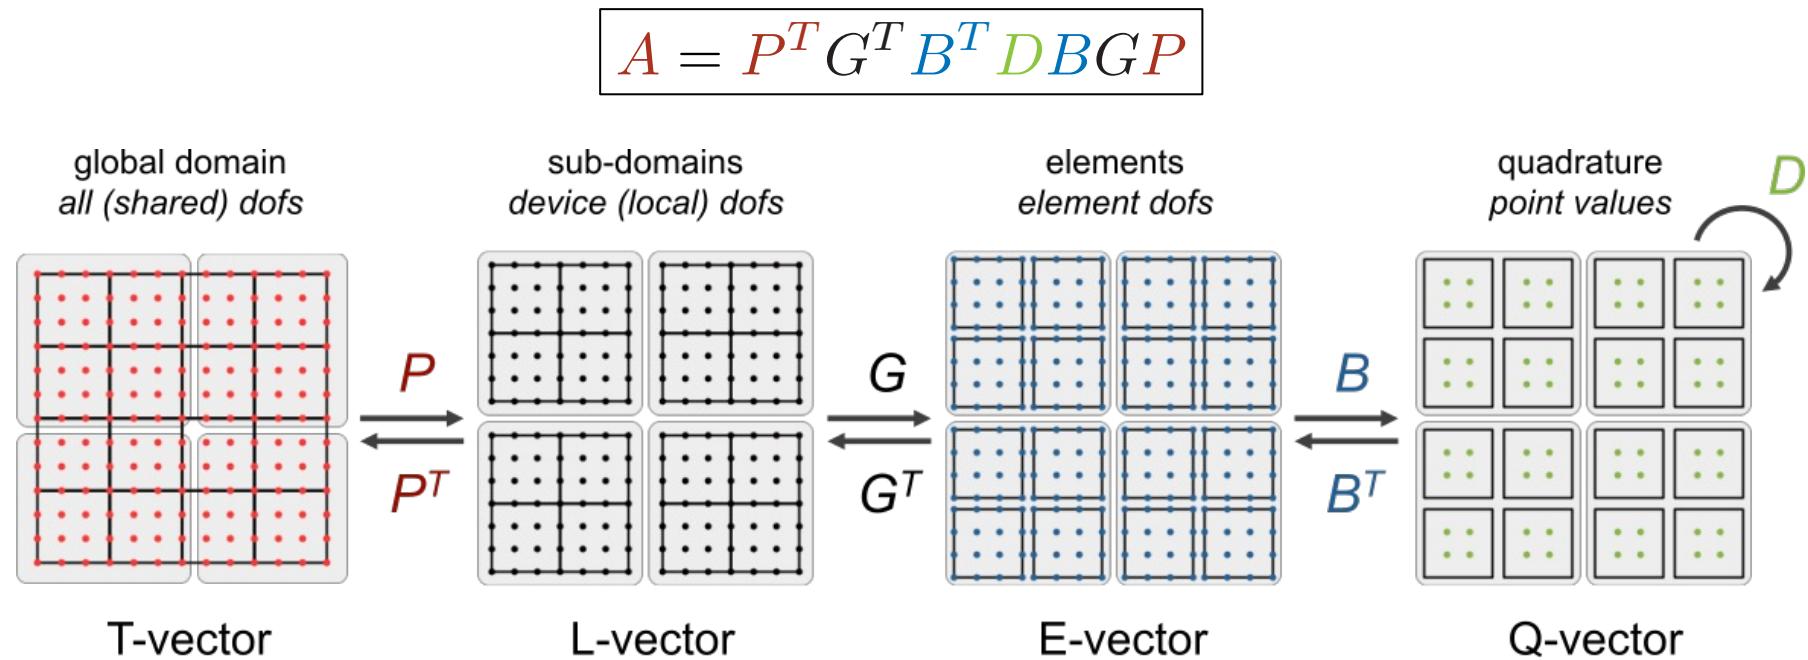


**Parallel partitioning via
Hilbert curve**

- weak+strong scaling up to ~400K MPI tasks on BG/Q
- **measure AMR only components:** interpolation matrix, assembly, marking, refinement & rebalancing (no linear solves, no “physics”)

Fundamental finite element operator decomposition

The assembly/evaluation of FEM operators can be decomposed into **parallel**, **mesh topology**, **basis**, and **geometry/physics** components:



- ✓ **partial assembly** = store only D , evaluate B (tensor-product structure)
- ✓ better representation than A : optimal memory, near-optimal FLOPs
- ✓ purely algebraic ✓ high-order operator format ✓ AD-friendly ∂_x



Example of a fast high-order operator

Poisson problem in variational form

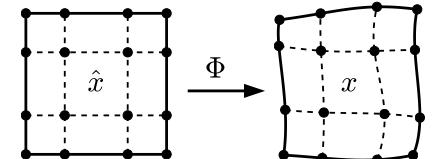
Find $u \in Q_p \subset \mathcal{H}_0^1$ s.t. $\forall v \in Q_p$,

$$\int_{\Omega} \nabla u \cdot \nabla v = \int_{\Omega} fv$$

Stiffness matrix (unit coefficient)

$$\begin{aligned} \int_{\Omega} \nabla \varphi_i \nabla \varphi_j &= \sum_E \int_E \nabla \varphi_i \nabla \varphi_j \\ &= \sum_E \sum_k \alpha_k J_E^{-1}(q_k) \hat{\nabla} \hat{\varphi}_i(q_k) J_E^{-1}(q_k) \hat{\nabla} \hat{\varphi}_j(q_k) |J_E(q_k)| \\ &= \sum_E \sum_k \underbrace{\hat{\nabla} \hat{\varphi}_i(q_k)}_{G, G^T} \underbrace{\left(\alpha_k J_E^{-T}(q_k) J_E^{-1}(q_k) |J_E(q_k)| \right)}_{D_{kk}} \underbrace{\hat{\nabla} \hat{\varphi}_j(q_k)}_{B_{kj}} \end{aligned}$$

- J is the Jacobian of the element mapping (geometric factors)

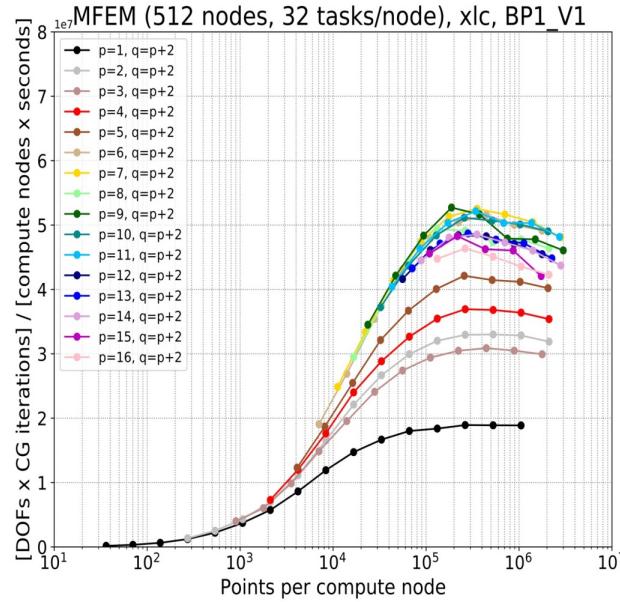


- G is usually Boolean (except AMR)
- Element matrices $A_E = B^T D B$, are full, account for bulk of the physics, can be applied in parallel

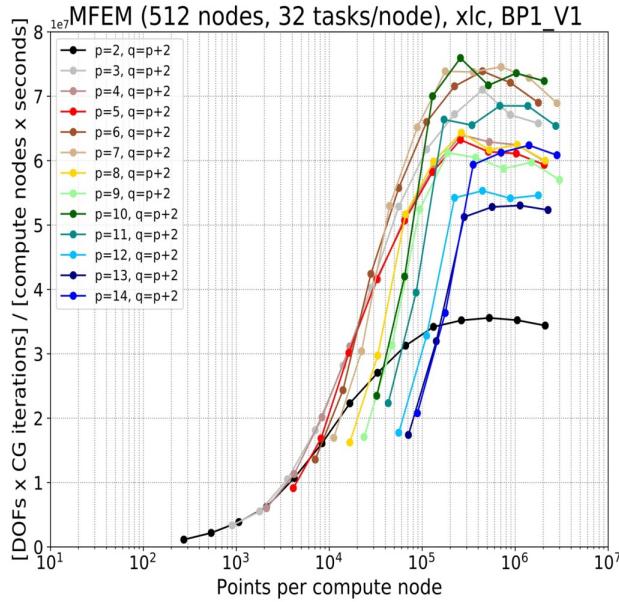
$$\begin{bmatrix} A^1 & & & \\ & A^2 & & \\ & & \ddots & \\ & & & A^4 \end{bmatrix}$$

- Never form A_E , just apply its action based on actions of B , B^T and D

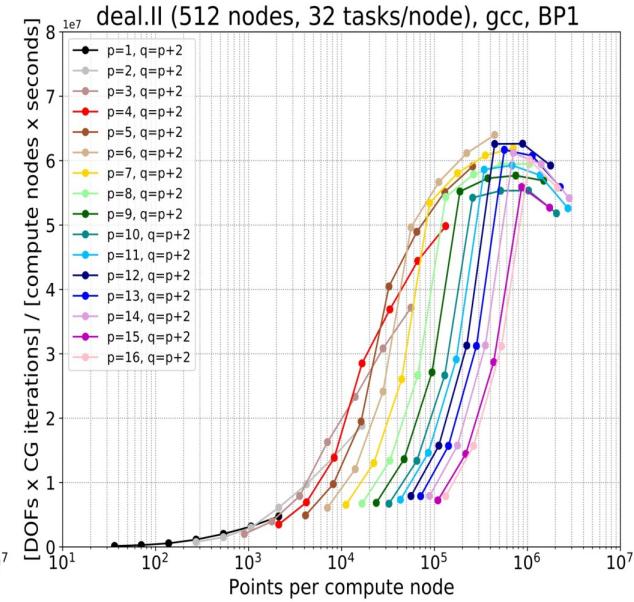
CEED BP1 bakeoff on BG/Q



Nek5000



MFEM-improved

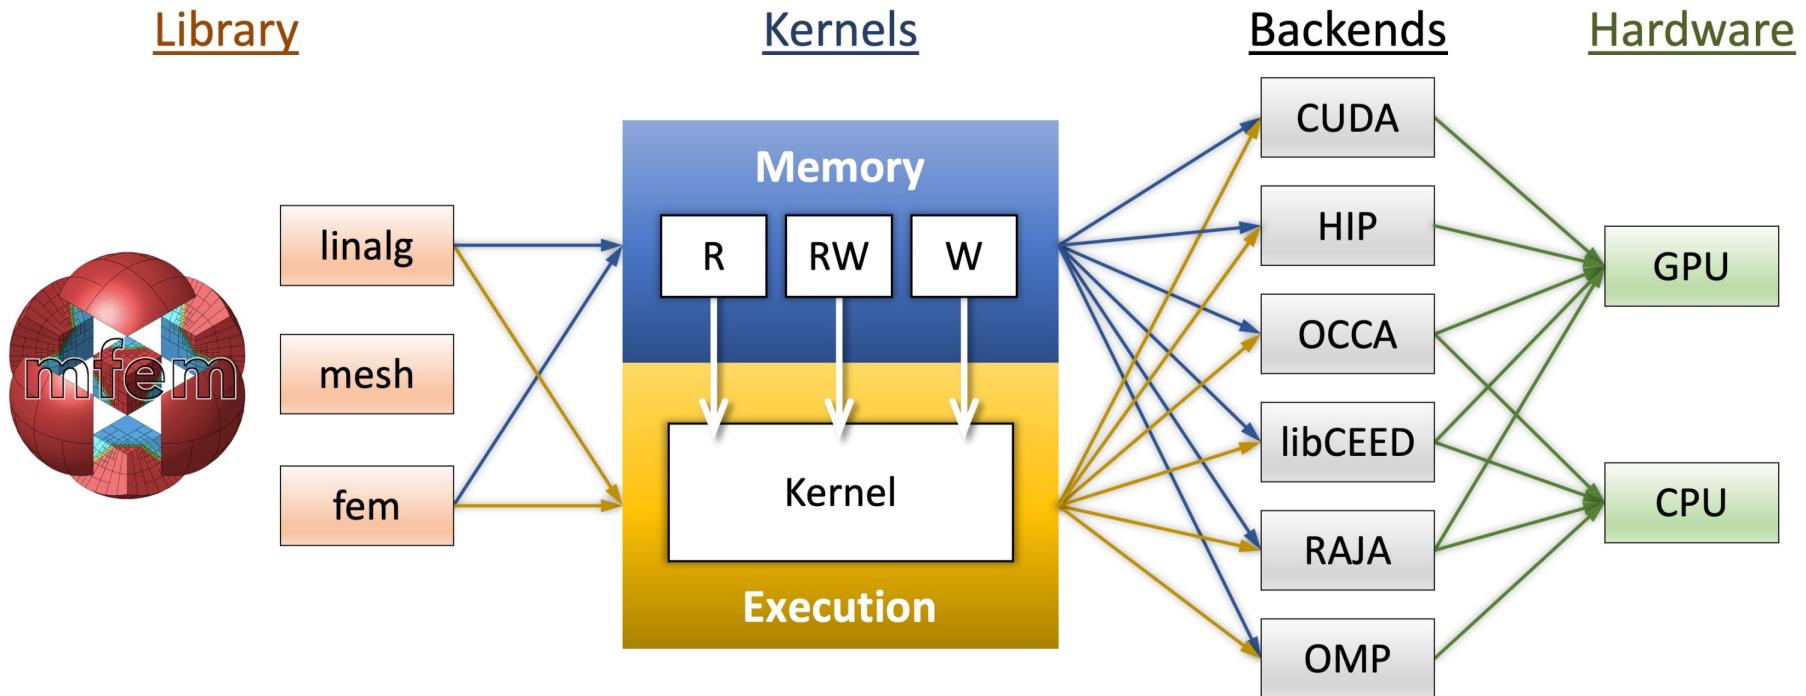


deal.ii

- ✓ All runs done on BG/Q (for repeatability), 16384 MPI ranks. Order $p = 1, \dots, 16$; quad. points $q = p + 2$.
- ✓ BP1 results of MFEM+xlc (left), MFEM+xlc+intrinsics (center), and deal.ii + gcc (right) on BG/Q.
- ✓ Paper: “Scalability of High-Performance PDE Solvers”, IJHPCA, 2020
- ✓ Cooperation/collaboration is what makes the bake-offs rewarding.

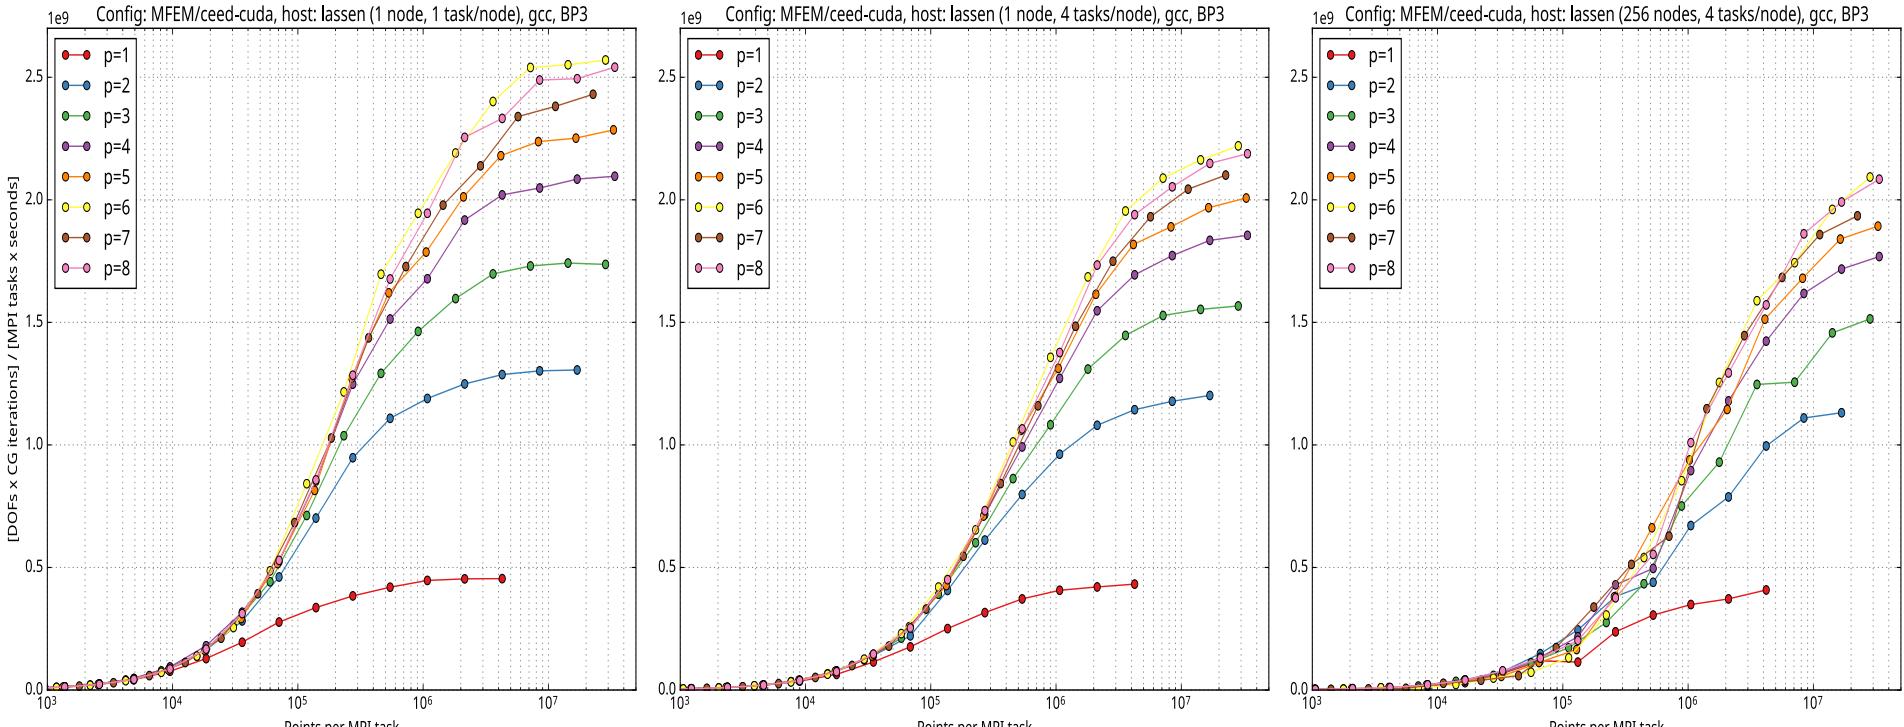
Device support in MFEM

MFEM supports GPU acceleration in many linear algebra and finite element operations



- Several MFEM examples + miniapps have been ported with small changes
- Many kernels have a single source for CUDA, RAJA and OpenMP backends
- Backends are runtime selectable, can be mixed
- Recent improvements in CUDA, HIP, RAJA, SYCL, ...

MFEM performance on multiple GPUs



1 GPU

4 GPUs

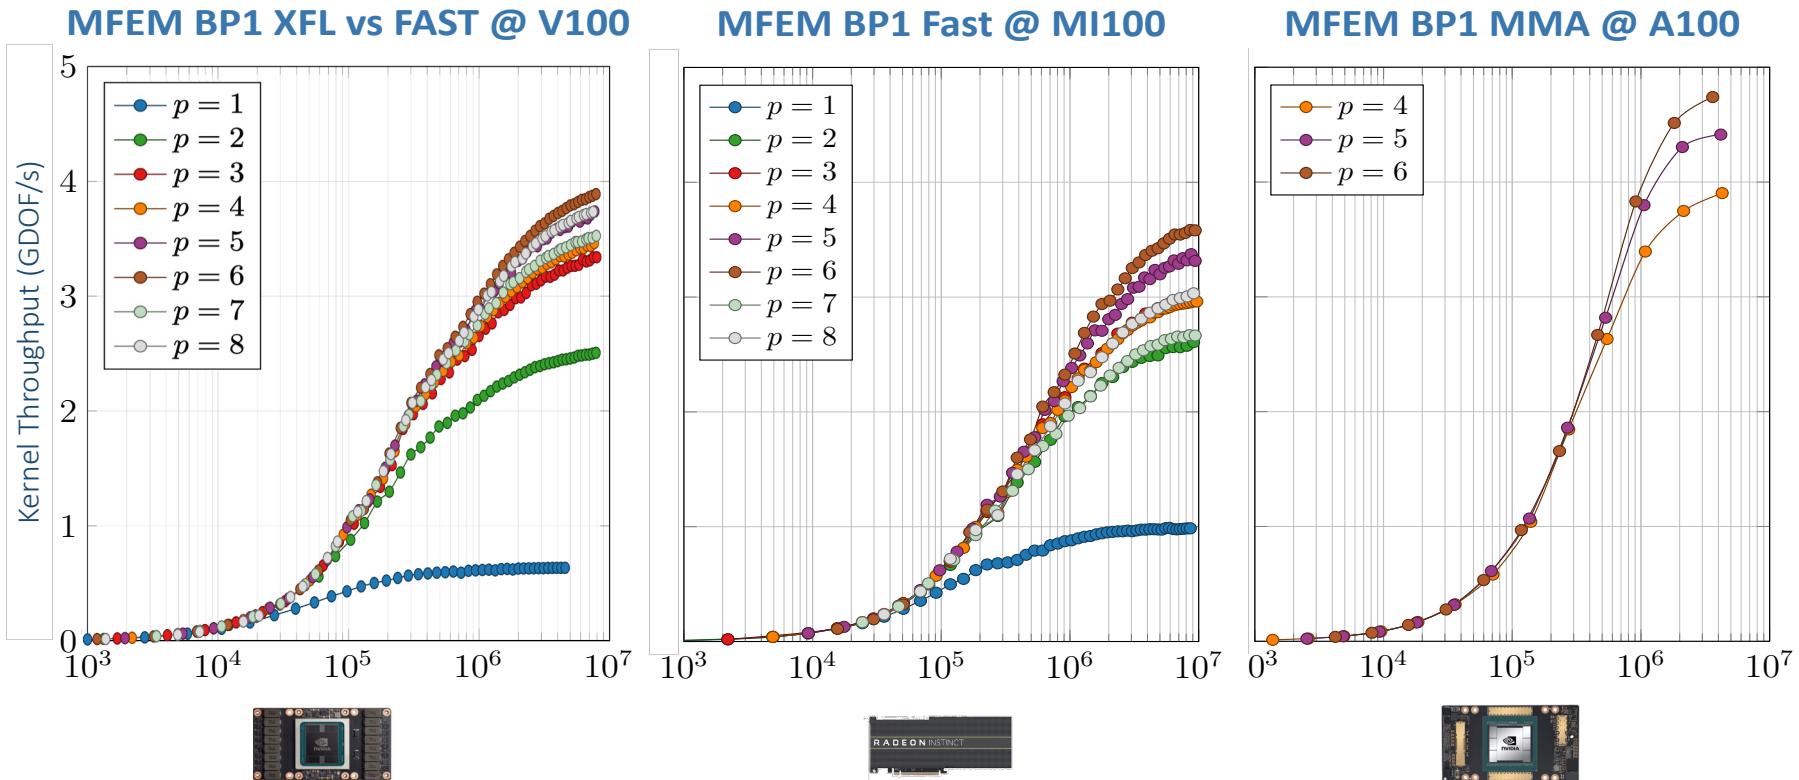
1024 GPUs

Single GPU performance: **2.6 GDOF/s**
Problem size: 10+ million

Best total performance: **2.1 TDOF/s**
Largest size: 34 billion

Optimized kernels for MPI buffer packing/unpacking on the GPU

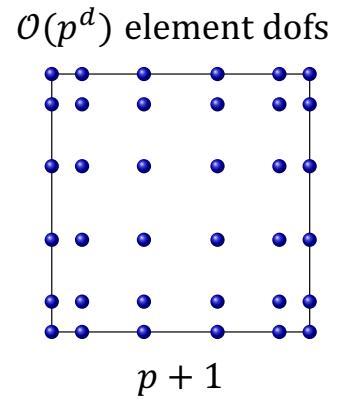
Recent improvements on NVIDIA and AMD GPUs



*New MFEM GPU kernels: perform on both V100 + MI100,
have better strong scaling,
can utilize tensor cores on A100*

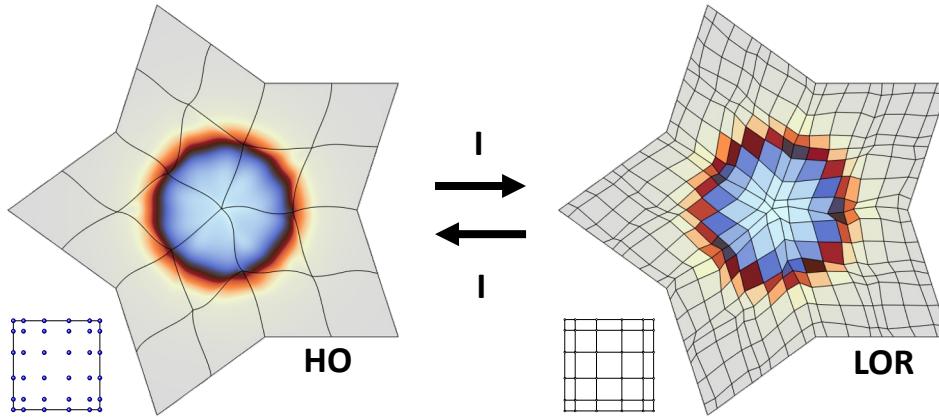
Matrix-free preconditioning

- **Explicit matrix assembly impractical at high order:**
 - Polynomial degree p , spatial dimension d
 - Matrix assembly + sparse matvecs:
 - $\mathcal{O}(p^{2d})$ memory transfers
 - $\mathcal{O}(p^{3d})$ computations
 - can be reduced to $\mathcal{O}(p^{2d+1})$ computations by sum factorization
 - Matrix-free action of the operator (partial assembly):
 - $\mathcal{O}(p^d)$ memory transfers – *optimal*
 - $\mathcal{O}(p^{d+1})$ computations – *nearly-optimal*
 - efficient iterative solvers ***if combined with effective preconditioners***
- **Challenges:**
 - Traditional matrix-based preconditioners (e.g. AMG) not available
 - Condition number of diffusion systems grows like $\mathcal{O}(p^3/h^2)$



Low-Order-Refined (LOR) preconditioning

Efficient LOR-based preconditioning of H_1 , $H(\text{curl})$, $H(\text{div})$ and L^2 high-order operators

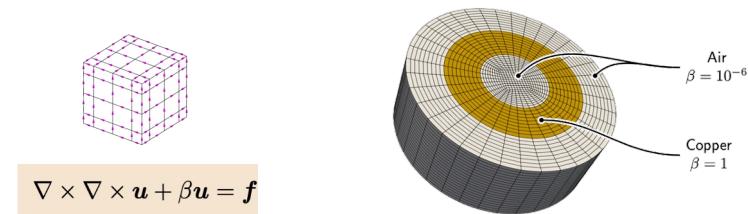


- Pick LOR space and HO basis so $\mathbf{P}=\mathbf{R}=\mathbf{I}$ (Gerritsma, Dohrmann)
- \mathbf{A}_{LOR} is sparse and spectrally equivalent to \mathbf{A}_{HO}

Theorem 2. Let M_\star and K_\star denote the mass and stiffness matrices, respectively, where \star represents one of the above-defined finite element spaces with basis as in Section 4.3. Then we have the following spectral equivalences, independent of mesh size h and polynomial degree p .

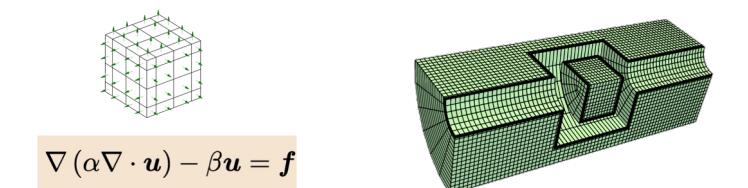
$$\begin{aligned} M_{V_h} &\sim M_{V_p}, & K_{V_h} &\sim K_{V_p}, \\ M_{W_h} &\sim M_{W_p}, & K_{W_h} &\sim K_{W_p}, \\ M_{X_h} &\sim M_{X_p}, & K_{X_h} &\sim K_{X_p}, \\ M_{Y_h} &\sim M_{Y_{p-1}}, \\ M_{Z_h} &\sim M_{Z_p}, & K_{Z_h} &\sim K_{Z_p}. \end{aligned}$$

- $(\mathbf{A}_{\text{HO}})^{-1} \approx (\mathbf{A}_{\text{LOR}})^{-1} \approx \mathbf{B}_{\text{LOR}}$ - can use BoomerAMG, AMS, ADS



LOR-AMS						
p	Its.	Assembly (s)	AMG Setup (s)	Solve (s)	# DOFs	# NNZ
2	41	0.082	0.277	0.768	516,820	1.65×10^7
3	63	0.251	0.512	2.754	1,731,408	5.64×10^7
4	75	0.679	1.133	7.304	4,088,888	1.34×10^8
5	62	1.574	2.185	11.783	7,968,340	2.61×10^8
6	89	3.336	4.024	30.702	13,748,844	4.51×10^8

Matrix-Based AMS						
p	Its.	Assembly (s)	AMG Setup (s)	Solve (s)	# DOFs	# NNZ
2	39	0.140	0.385	1.423	516,820	5.24×10^7
3	44	1.368	1.572	9.723	1,731,408	4.01×10^8
4	49	9.668	5.824	45.277	4,088,888	1.80×10^9
5	53	61.726	15.695	148.757	7,968,340	5.92×10^9
6	56	502.607	40.128	424.100	13,748,844	1.59×10^{10}



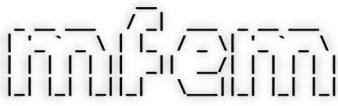
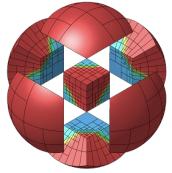
p	LOR-ADS		Matrix-Based ADS		Speedup
	Runtime (s)	Memory (GB)	Runtime (s)	Memory (GB)	
2	2.11	0.04	2.98	0.20	1.41×
3	6.64	0.15	22.58	1.84	3.40×
4	17.40	0.35	114.35	9.13	6.57×
5	43.70	0.68	422.74	32.21	9.67×
6	92.76	1.18	1324.94	91.09	14.28×

High-order methods show promise for high-quality & performance simulations on exascale platforms

- **More information and publications**

- MFEM – mfem.org
- BLAST – computation.llnl.gov/projects/blast
- CEED – ceedexascaleproject.org

- **Open-source software**



CEED
EXASCALE DISCRETIZATIONS

- **Ongoing R&D**

- GPU-oriented algorithms for Frontier, Aurora, El Capitan
- Matrix-free scalable preconditioners
- Automatic differentiation, design optimization
- Deterministic transport, multi-physics coupling



Q4 Rayleigh-Taylor single-material ALE on 256 processors

Upcoming MFEM Events

MFEM in the Cloud Tutorial

August 10, 2023

<https://mfem.org/tutorial>

MFEM Community Workshop

October 26, 2023



<https://mfem.org/workshop>



Seminar series: <https://mfem.org/seminar>



**This work performed under the auspices of the U.S. Department of Energy by
Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344.
LLNL-PRES-755924**

Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

Unstructured Mesh Methods

Unstructured mesh – a spatial domain discretization composed of topological entities with general connectivity and shape

Advantages

- Automatic mesh generation for any level of geometric complexity
- Can provide the highest accuracy on a per degree of freedom basis
- General mesh anisotropy possible
- Meshes can easily be adaptively modified
- Given a complete geometry, with analysis attributes defined on that model, the entire simulation workflow can be automated

Disadvantages

- More complex data structures and increased program complexity, particularly in parallel
- Requires careful mesh quality control (level of control required is a function of the unstructured mesh analysis code)
- Poorly shaped elements increase condition number of global system – makes matrix solves harder
- Non-tensor product elements not as computationally efficient

Unstructured Mesh Methods

Goal of FASTMath unstructured mesh developments include:

- Provide unstructured mesh components that are easily used by application code developers to extend their simulation capabilities
- Ensure those components execute on exascale computing systems and support performant exascale application codes
- Develop components to operate through multi-level APIs that increase interoperability and ease of integration
- Address technical gaps by developing tools that address needs and/or eliminate/minimize disadvantages of unstructured meshes
- Work with DOE application developers on integration of these components into their codes

FASTMath Unstructured Mesh Developments

Technology development areas:

- Unstructured Mesh Analysis Codes – Support application's PDE solution needs – MFEM library is a key example
- Performant Mesh Adaptation – Parallel mesh adaptation to integrate into analysis codes to ensure solution accuracy
- Dynamic Load Balancing and Task Management – Technologies to ensure load balance and effectively execute by optimal task placement
- Unstructured Mesh for Particle In Cell (PIC) Codes – Tools to support PIC on unstructured meshes
- Unstructured Mesh ML and UQ – ML for data reduction, adaptive mesh UQ
- In Situ Vis and Data Analytics – Tools to gain insight as simulations execute

FASTMath Unstructured Mesh Tools and Components

- FE Analysis codes
 - MFEM (<https://mfem.org/>)
 - LGR (<https://github.com/SNLComputation/lgrtk>)
 - PHASTA (<https://github.com/phasta/phasta>)
- Unstructured Mesh Infrastructure
 - Omega_h (https://github.com/SNLComputation/omega_h)
 - PUMI/MeshAdapt (<https://github.com/SCOREC/core>)
 - PUMIpic (<https://github.com/SCOREC/pumi-pic>)
- Load balancing, task placement
 - Zoltan (<https://github.com/sandialabs/Zoltan>)
 - Zoltan2 (<https://github.com/trilinos/Trilinos/tree/master/packages/zoltan2>)
 - Xtra-PULP (<https://github.com/HPCGraphAnalysis/PuLP>)
 - EnGPar (<http://scorec.github.io/EnGPar/>)
- Unstructured Mesh PIC applications
 - XGCM (<https://github.com/SCOREC/xgcm>)
 - GITRm (<https://github.com/SCOREC/gitrm>)

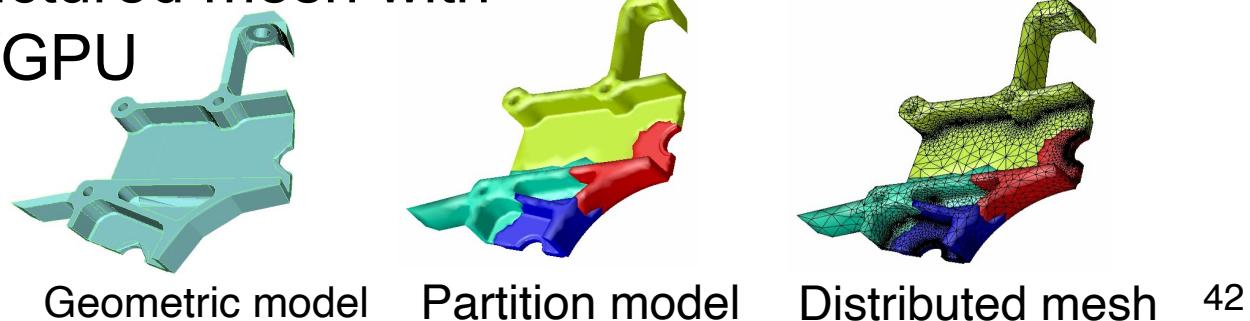
Parallel Unstructured Mesh Infrastructure

Support unstructured mesh interactions on exascale systems

- Mesh hierarchy to support interrogation and modification
- Maintains linkage to original geometry
- Conforming mesh adaptation
- Inter-process communication
- Supports field operations

Tools

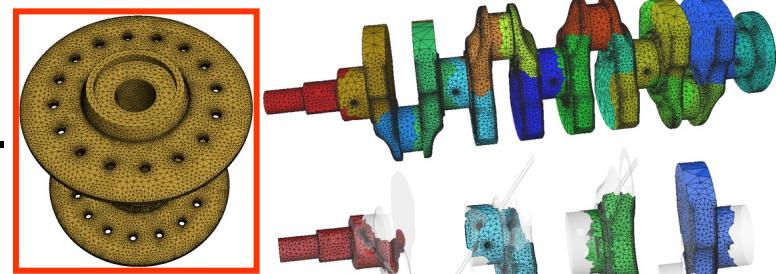
- Omega_h – CPU/GPU support
- PUMI – CPU based curved mesh adapt.
- PUMIPic – Unstructured mesh with particles for CPU/GPU



Mesh Generation and Control

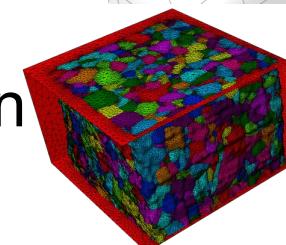
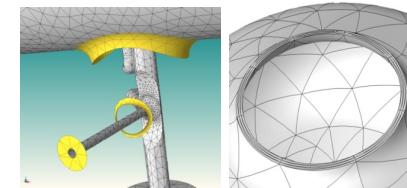
Mesh Generation:

- Automatically mesh complex domains – should work directly from CAD, image data, etc.
- Use tools like Gmsh, Simmetrix, etc.



Mesh control:

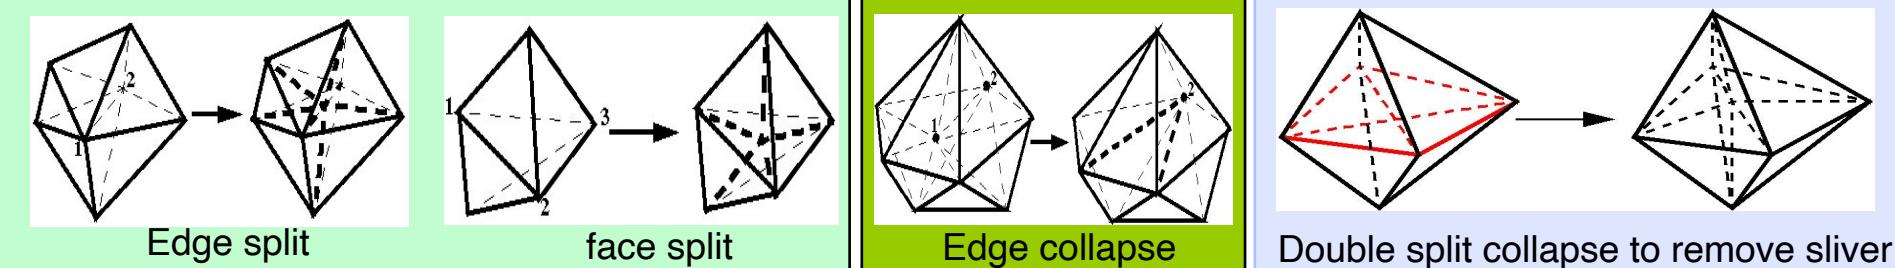
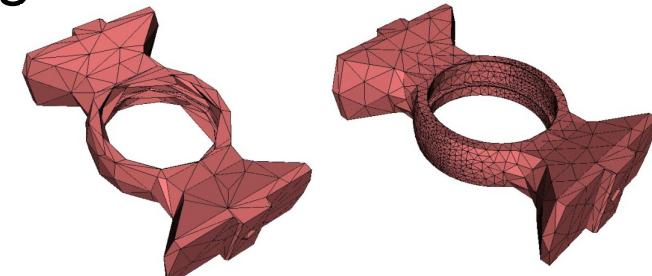
- Use *a posteriori* information to improve mesh
- Curved geometry and curved mesh entities
- Support full range of mesh modifications – vertex motion, mesh entity curving, cavity based refinement and coarsening, etc. anisotropic adaptation
- Control element shapes as needed by the various discretization methods for maintaining accuracy and efficiency



Parallel execution of all functions is critical on large meshes

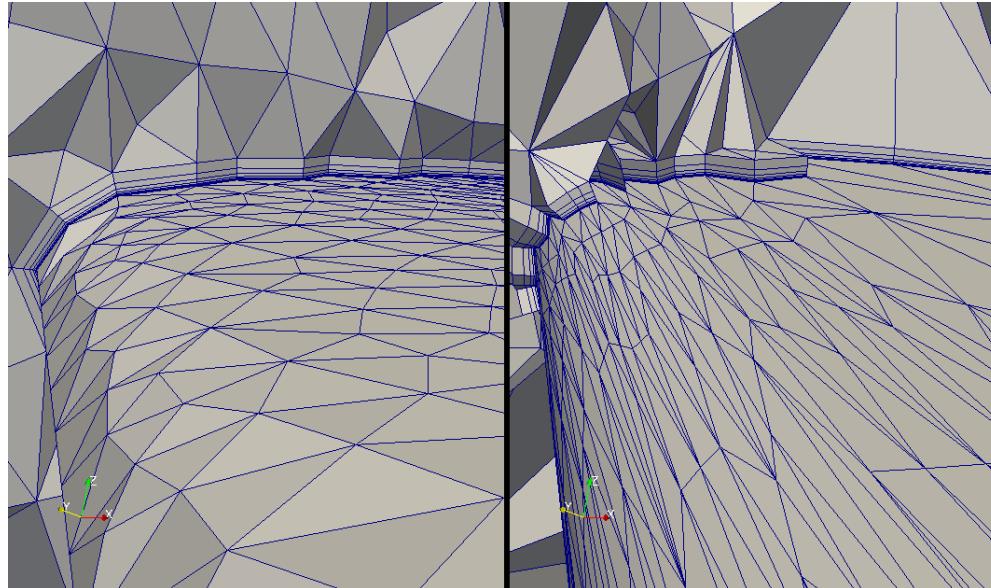
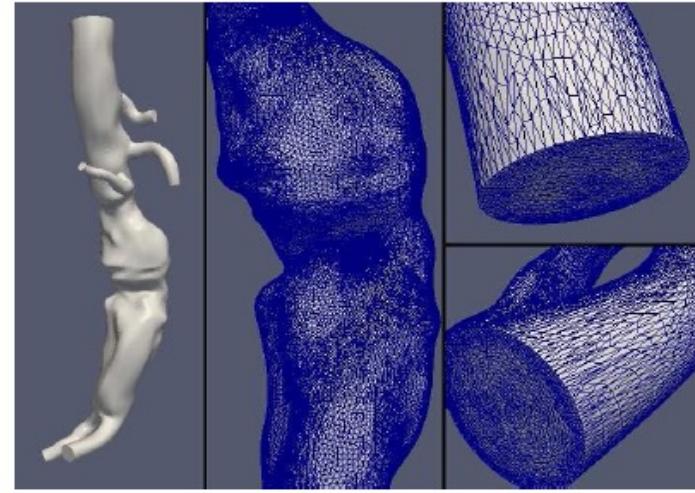
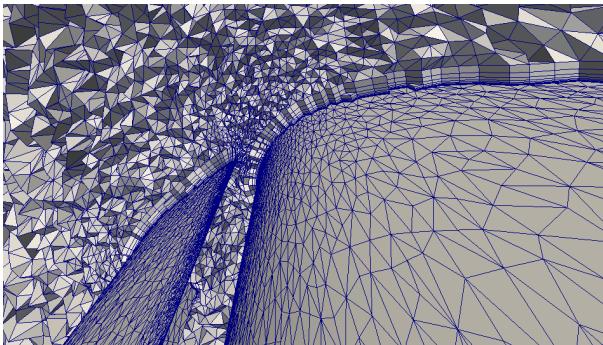
General Mesh Modification for Mesh Adaptation

- Driven by an anisotropic mesh size field that can be set by any combination of criteria
- Employ a “complete set” of mesh modification operations to alter the mesh into one that matches the given mesh size field
- Advantages
 - Supports general anisotropic meshes
 - Can obtain level of accuracy desired
 - Can deal with any level of geometric domain complexity
 - Solution transfer can be applied incrementally - provides more control to satisfy conservation constraints



Mesh Adaptation Status

- Applied to very large scale models
 - 92B elements on 3.1M processes on $\frac{3}{4}$ million cores
- Local solution transfer supported through callback
- Effective storage of solution fields on meshes
- Supports adaptation with boundary layer meshes

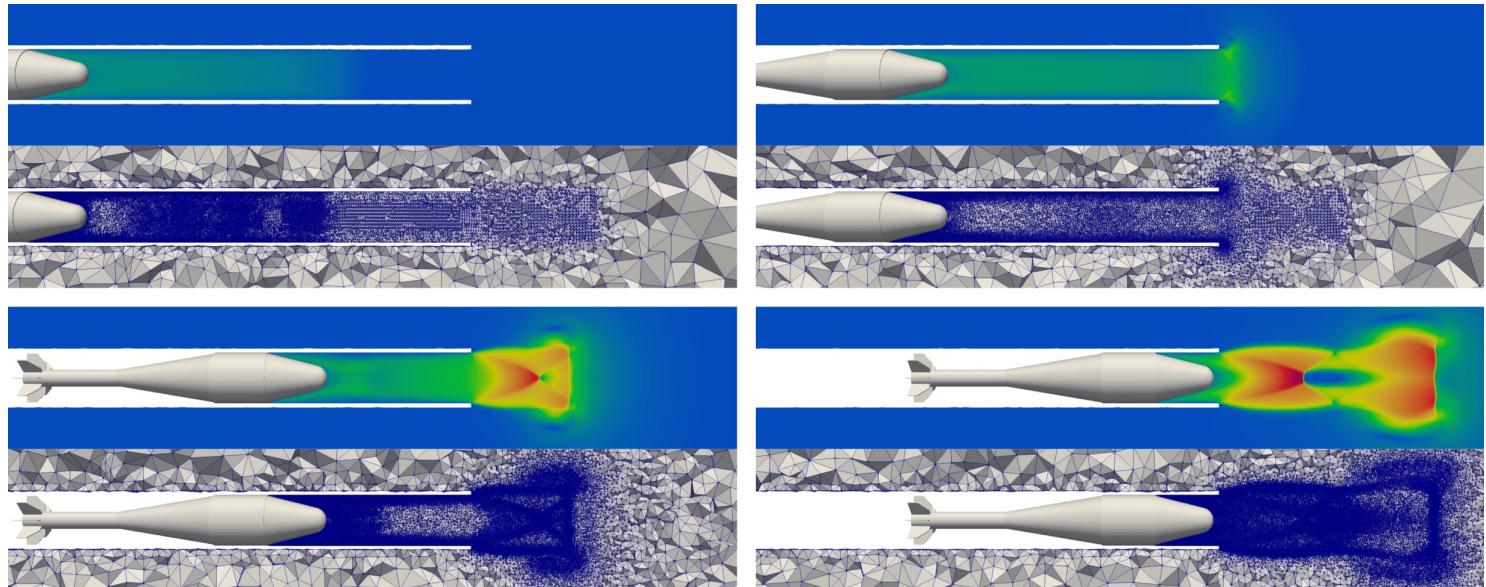


Mesh Adaptation of Evolving Geometry Problems

Many applications have geometry that evolves as a function of the results –
Effective adaptive loops combine mesh motion and mesh modification

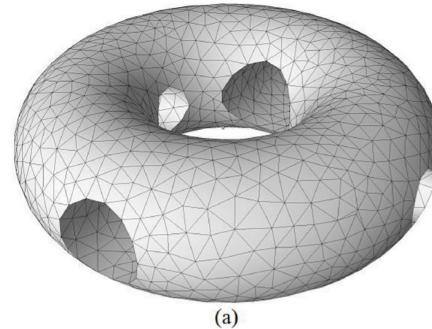
Adaptive loop:

1. Initialize analysis case, generate initial mesh, start time stepping loop
2. Perform time steps employing mesh motion - monitor element quality and discretization errors
3. When element quality is not satisfactory or discretization errors too large – set mesh size field and perform mesh modification
4. Return to step 2.

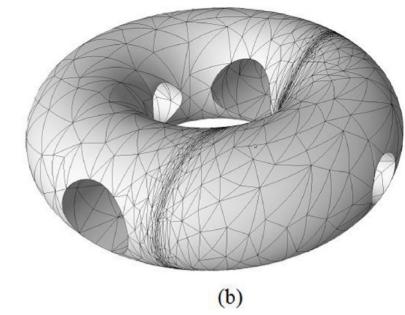


Mesh Adaptation

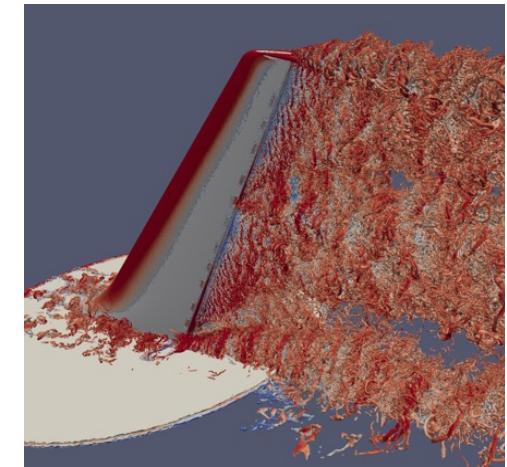
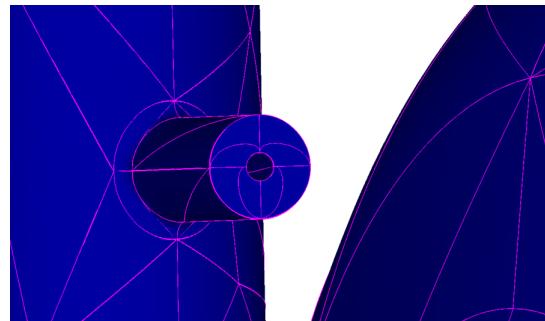
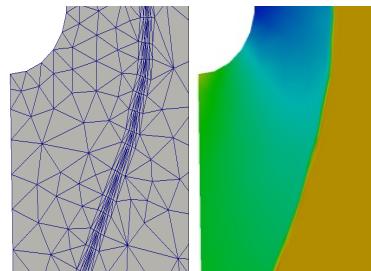
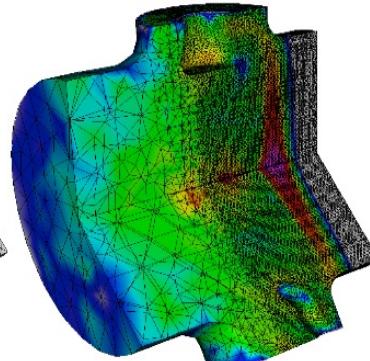
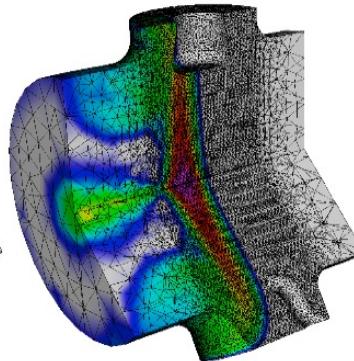
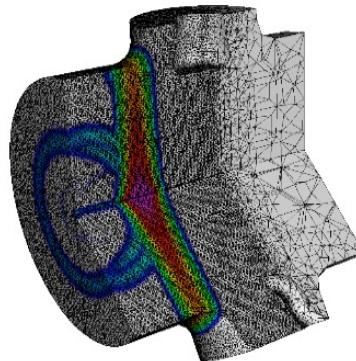
- Supports adaptation of curved elements
- Adaptation based on multiple criteria, examples
 - Level sets at interfaces
 - Tracking particles
 - Discretization errors
 - Controlling element shape in evolving geometry



(a)

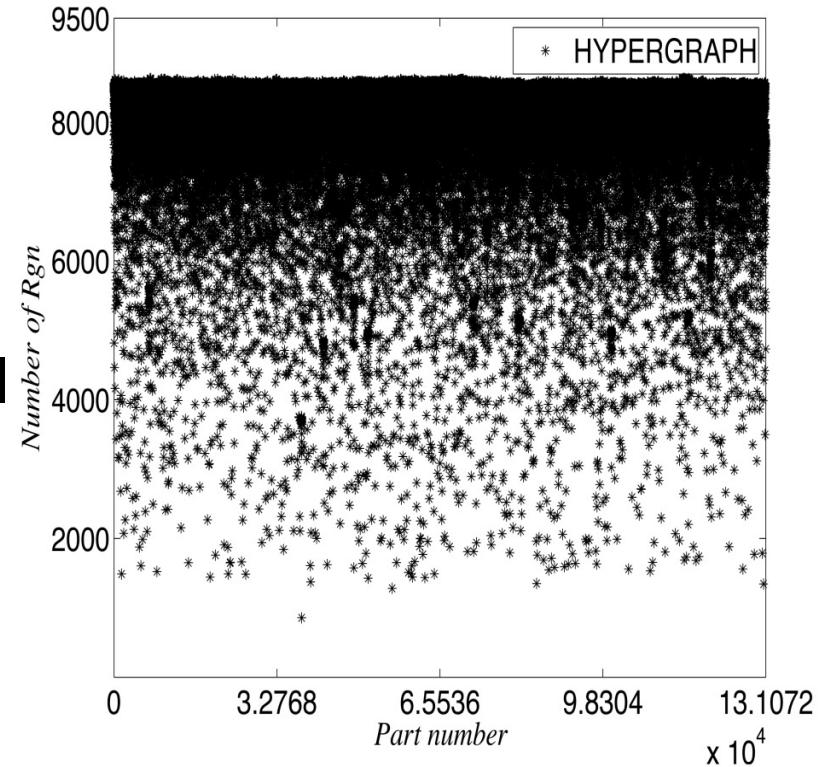


(b)



Load Balancing, Dynamic Load balancing

- Purpose: Balance or rebalance computational load while controlling communications
 - Equal “work load” with minimum inter-process communications
- FASTMath load balancing tools
 - Zoltan/Zoltan2 libraries provide multiple dynamic partitioners with general control of partition objects and weights
 - EnGPar diffusive multi-criteria partition improvement
 - XtraPuLP multi-constraint multi-objective label propagation-based graph partitioner



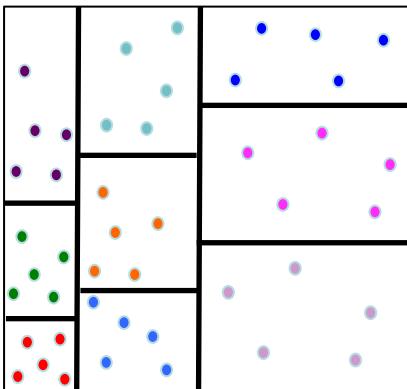
Architecture-aware partitioning and task mapping

- *Reduce application communication time at extreme scale*
- *Partitioning and load balancing*: assign work to processes in ways that avoid process idle time and minimize communication
- *Task mapping*: assign processes to cores/GPUs in ways that reduce messages distances and network congestion
- *Important in extreme-scale systems*:
 - Small load imbalances can waste many resources
 - Large-scale networks can cause messages to travel long routes and induce congestion
- *Algorithms developed to*:
 - Account for underlying architectures & hierarchies
 - Run effectively side-by-side with application across many platforms (multicore, GPU)

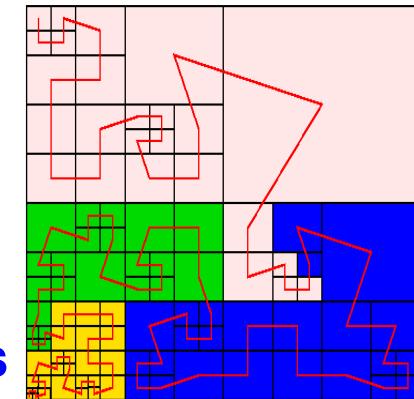
Zoltan/Zoltan2 Toolkits: Partitioners

*Suite of partitioners supports a wide range of applications;
no single partitioner is best for all applications.*

Geometric

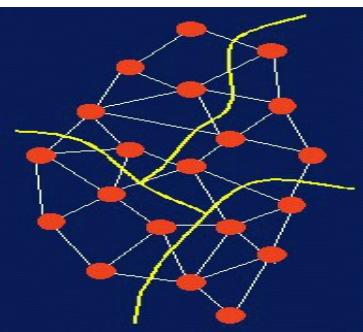


Recursive Coordinate Bisection
Recursive Inertial Bisection
Multi-Jagged Multi-section

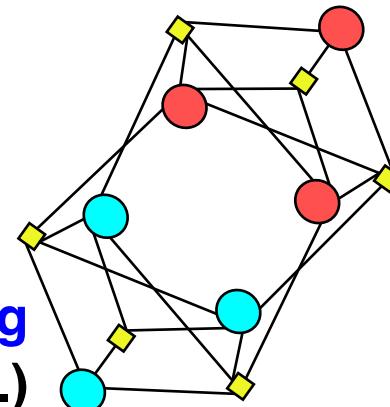


Space Filling Curves

Topology-based



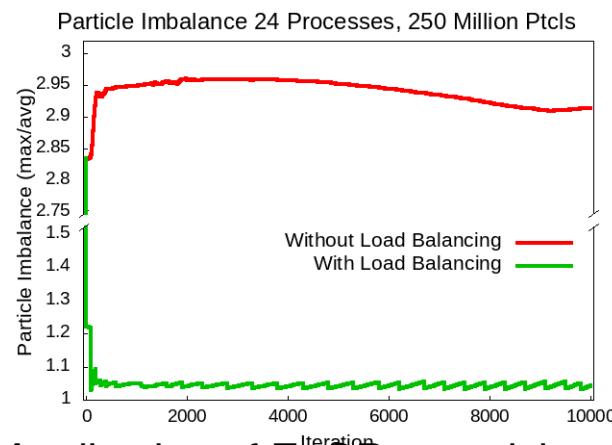
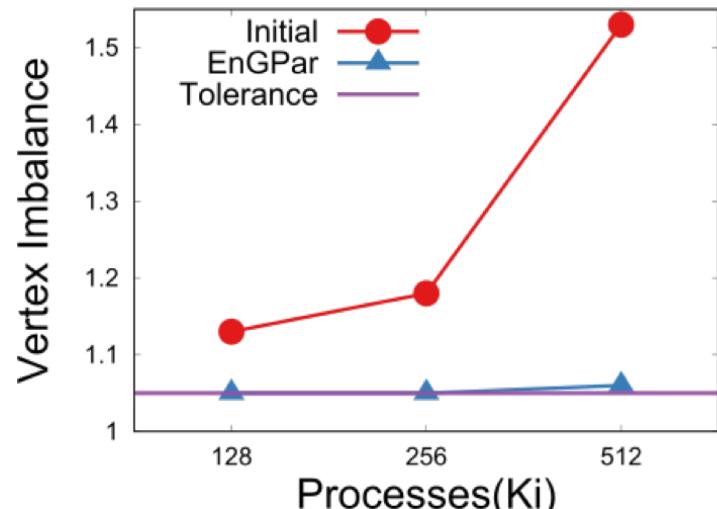
PHG Graph Partitioning
Interface to ParMETIS (U. Minnesota)
Interface to PT-Scotch (U. Bordeaux)



PHG Hypergraph Partitioning
Interface to PaToH (Ohio St.)

EnGPar quickly reduces large imbalances on (hyper)graphs with billions of edges on up to 512K processes

- Multi-(hyper)graph supports multiple dependencies (edges) between application work/data items (vertices)
- Application defined vertex and edges
- Diffusion sending if work from heavily loaded parts to lightly loaded parts
- In 8 seconds, EnGPar reduced a 53% vtx imbalance to 6%, at a cost of 5% elm imbalance, and edge cut increase by 1% on a 1.3B element mesh
- Applied to PIC calculations to support particle balance – 20% reduction in total run time

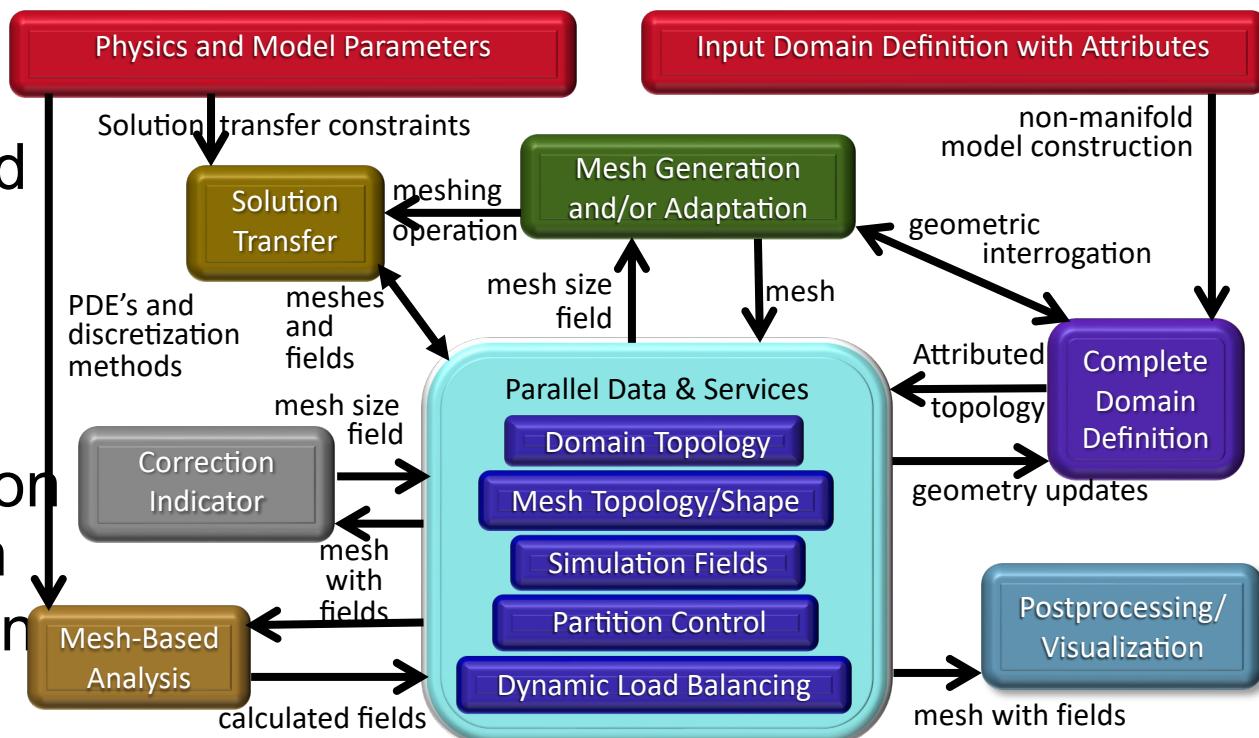


Application of EnGPar particle dynamic load balancing in a GITrm impurity transport simulation

Creation of Parallel Adaptive Loops

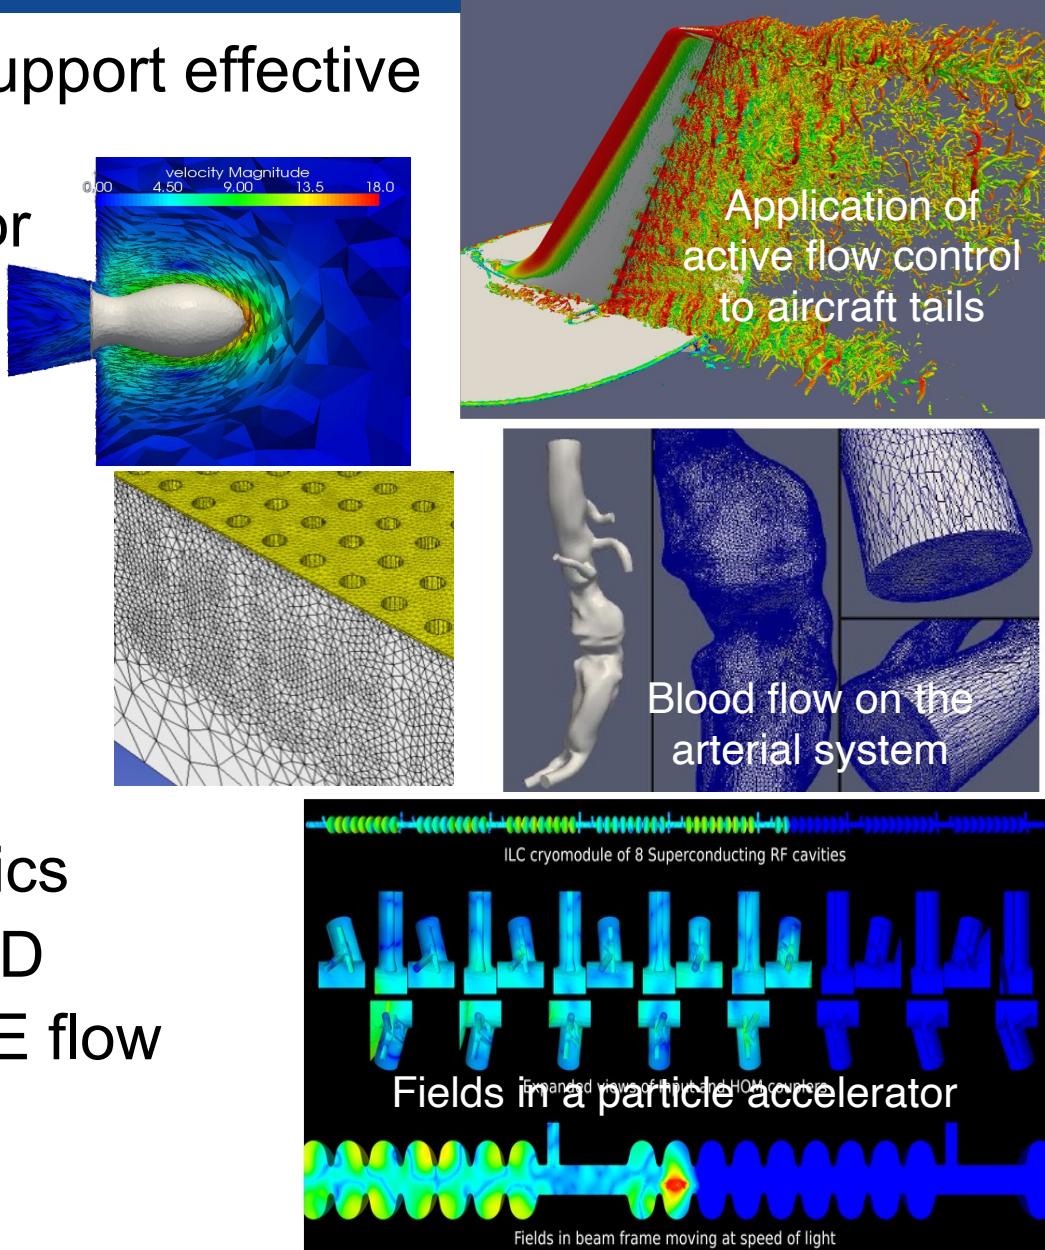
Parallel data and services used to develop adaptive loop

- Geometric model topology for domain linkage
- Mesh topology – it must be distributed
- Simulation fields distributed over geometric model and mesh
- Partition control
- Dynamic load balancing required at multiple steps
- API's to link to
 - CAD
 - Mesh generation and adaptation
 - Error estimation
 - etc.



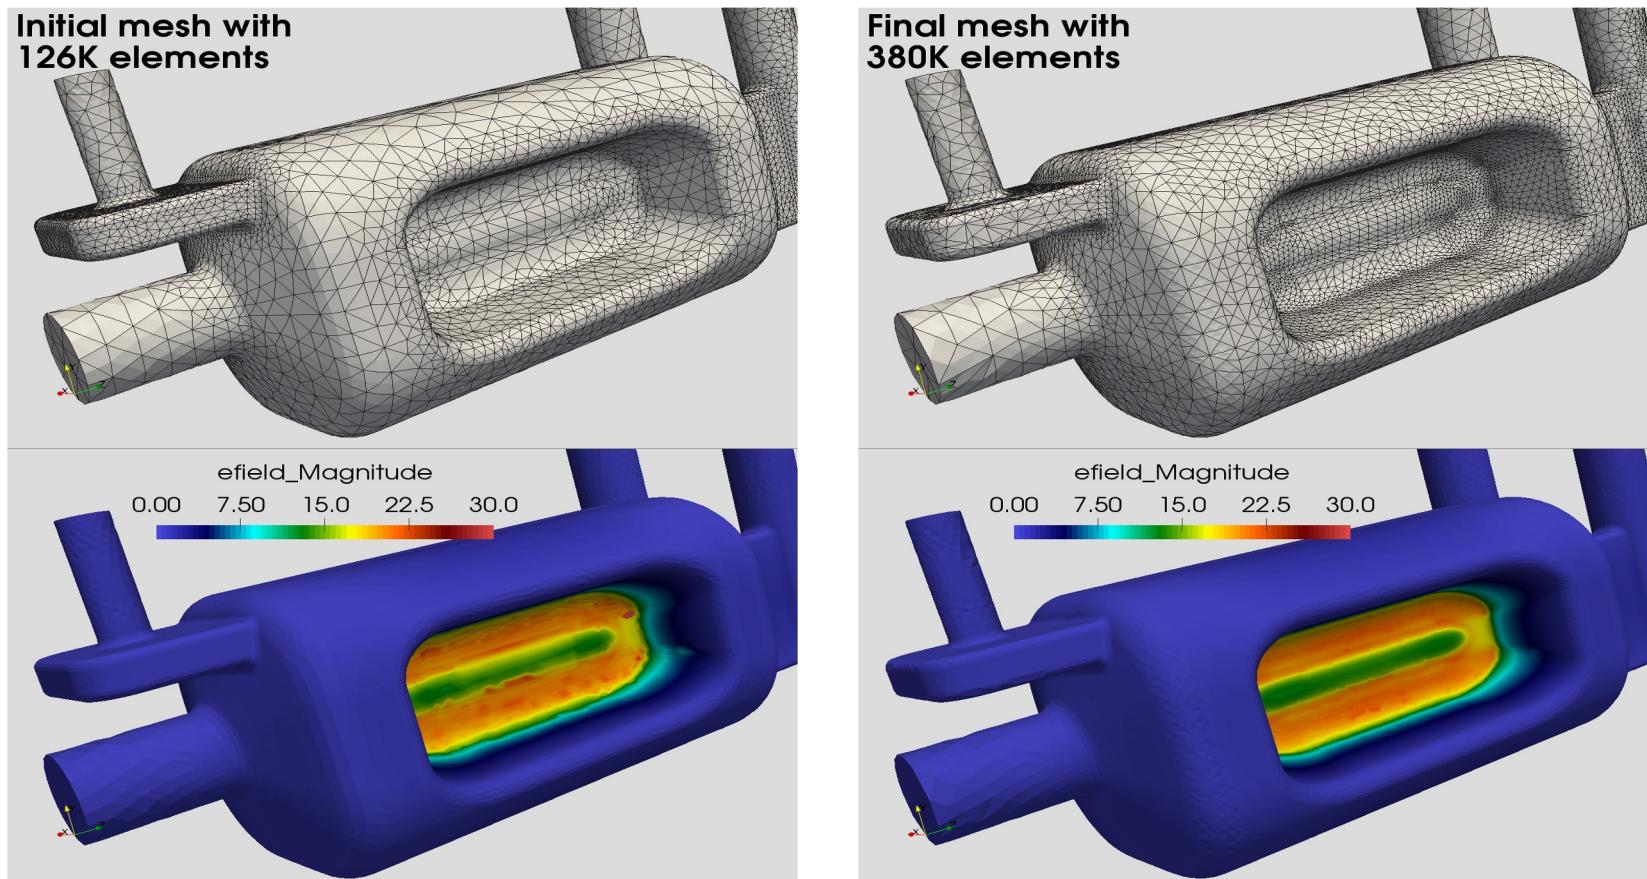
Parallel Adaptive Simulation Workflows

- In memory adaptive loops support effective data movement
- In-memory adaptive loops for
 - MFEM – High order FE framework
 - PHASTA – FE for NS
 - FUN3D – FV CFD
 - Proteus – multiphase FE
 - Albany – FE framework
 - ACE3P – High order FE electromagnetics
 - M3D-C1 – FE based MHD
 - Nektar++ – High order FE flow



Application interactions – Accelerator EM

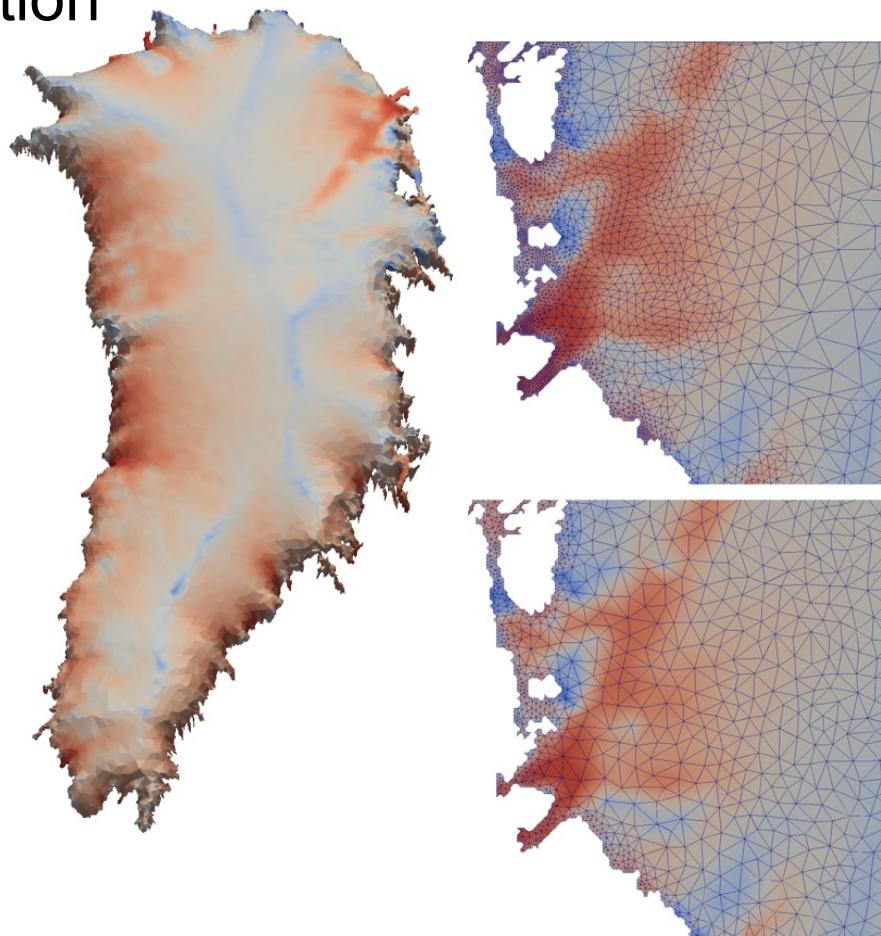
Omega3P Electro Magnetic Solver (second-order curved meshes)



This figure shows the adaptation results for the CAV17 model. (top left) shows the initial mesh with ~126K elements, (top right) shows the final (after 3 adaptation levels) mesh with ~380K elements, (bottom left) shows the first eigenmode for the electric field on the initial mesh, and (bottom right) shows the first eigenmode of the electric field on the final (adapted) mesh.

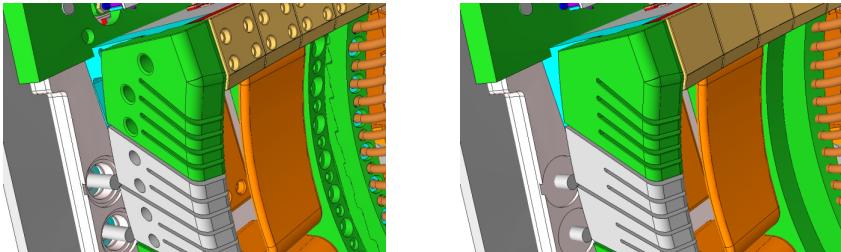
Application interactions – Land Ice

- FELIX, a component of the Albany framework is the analysis code
- Omega_h parallel mesh adaptation is integrated with Albany to do:
 - Estimate error
 - Adapt the mesh
- Ice sheet mesh is modified to minimize degrees of freedom
- Field of interest is the ice sheet velocity

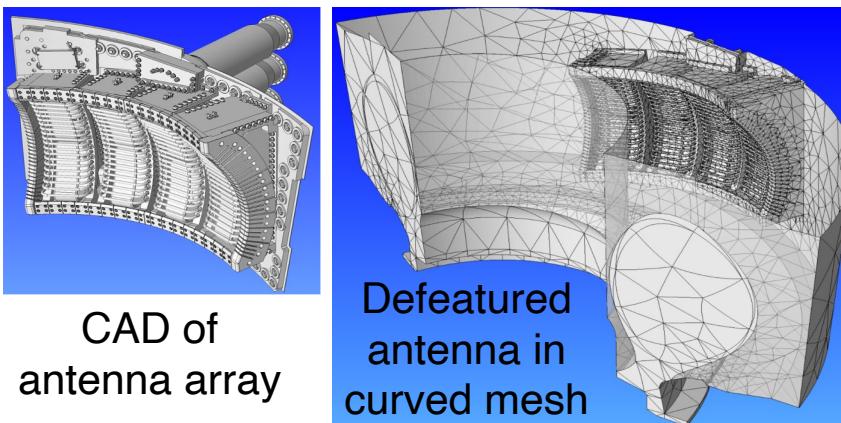


Application interactions – RF Fusion

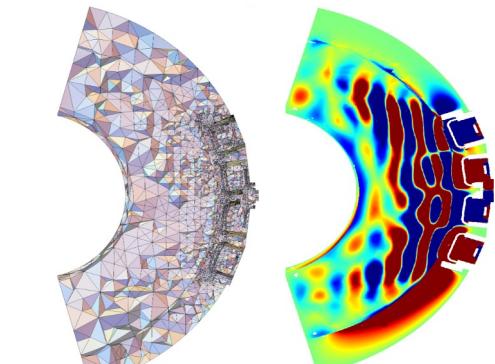
- Accurate RF simulations require
 - Detailed antenna CAD geometry
 - CAD geometry defeaturing
 - Extracted physics curves from GEQDSK equilibrium file
 - Analysis geometry combines CAD, and physics geometry
 - 3D meshes for accurate FE calculations in MFEM
 - Projection based error estimator
 - Conforming mesh adaptation with PUMI



Fast elimination of unwanted features



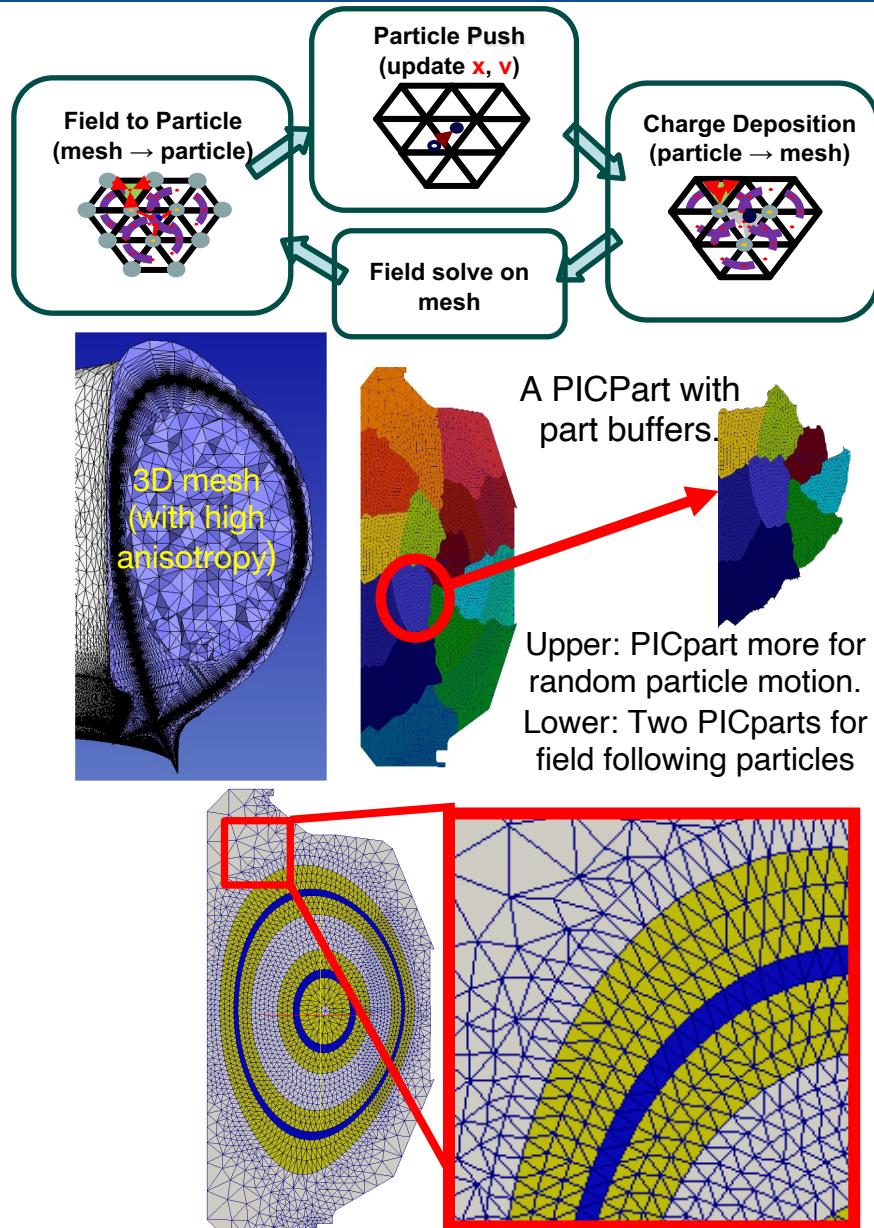
CAD of antenna array



Supporting Unstructured Mesh for Particle-in-Cell Calculations

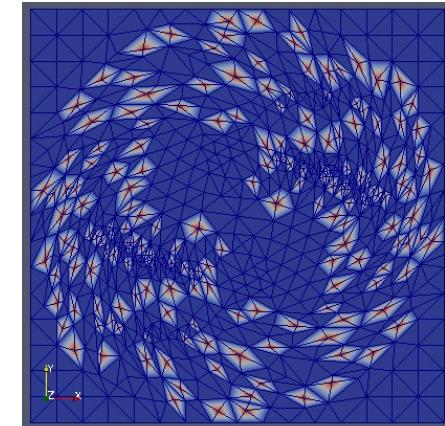
PUMIPic data structures are mesh centric

- Mesh is distributed as needed by the application in terms of PICparts
- Mesh can be graded and anisotropic
- Particle data associated with elements
- Operations take advantage of distributed mesh topology
- Mesh distributed in PICparts
 - Start with a partition of mesh into a set of “core parts”
 - A PICpart is defined by a “core part” and sufficient buffer to keep particles on process for one or more pushes
 - GPU version defines buffer as set of neighboring parts

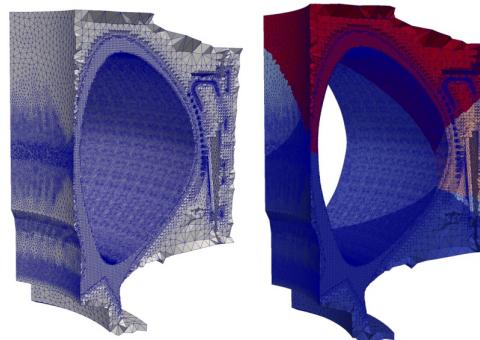


Mesh Data Structure for Heterogeneous Systems

- Mesh topology/adaptation tool - Omega
 - Conforming mesh adaptation (coarsening past initial mesh, refinement, swap)
 - Manycore and GPU parallelism using Kokkos
 - Distributed mesh via mesh partitions with MPI communications
 - Support for mesh-based fields
- Recent developments:
 - Curved mesh adaptation
 - More efficient field storage
 - Kokkos implementation on latest NVIDIA, AMD and Intel GPUs



Adaptation following
rotating flow field.



Serial and RIB partitioned mesh of RF
antenna and vessel model.

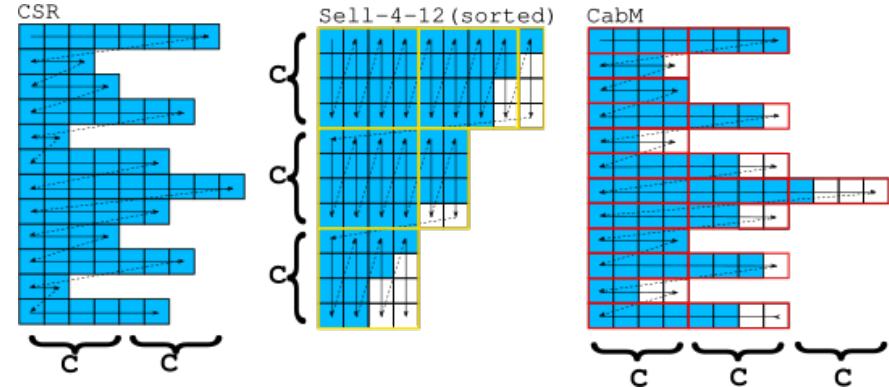
triangle	0	1
adj vertex	0 1 3 2 3 1	

adj triangle	0 0 1 1 0 1	-1
offset	0 1 3 4 6	

Mesh entity adjacency arrays.

PUMIPic Particle Data Structures

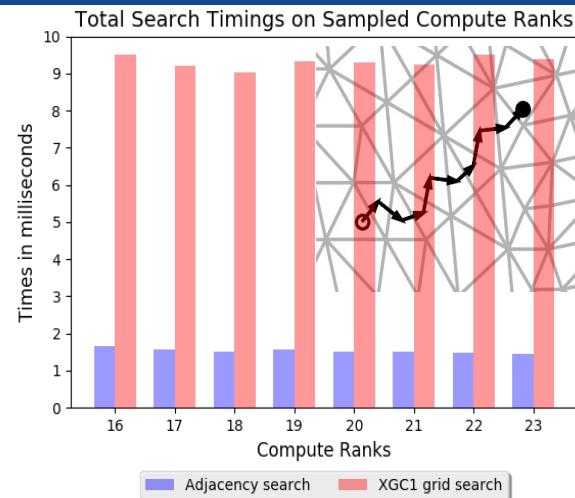
- Layout of particles in memory is critical to performance
 - Optimizes push (sort/rebuild), scatter, and gather operations
 - Associate particles with elements for large per element particle cases
 - Support changes in the number of particles per element
 - Evenly distributes work under a range of particle distributions (e.g. uniform, Gaussian, exponential, etc.)
 - Stores a lot of particles per GPU – low overhead
- Particle data structure interface and implementation
 - API abstracts implementation for PIC code developers
 - CSR, Sell-C- σ , CabanaM
 - Performance is a function of particle distribution
 - Cabana AoSoA w/a CSR index of elements-to-particles are promising



Left to Right: CSR, SCS with vertical slicing (yellow boxes), CabanaM (red boxes are SOAs). C is a team of threads.

PIC Operations Supported by PUMIPic

- Particle push
- Adjacency based search
 - Faster than grid based search
- Element-to-particle association update
- Particle Migration
- Particle path visualization
- Mesh partitioning w/buffer regions
- Mesh field association
- Poisson field solve using PETSc DMFlex on GPUs
- Checkpoint/restart of particle and mesh data – supports customization for each application



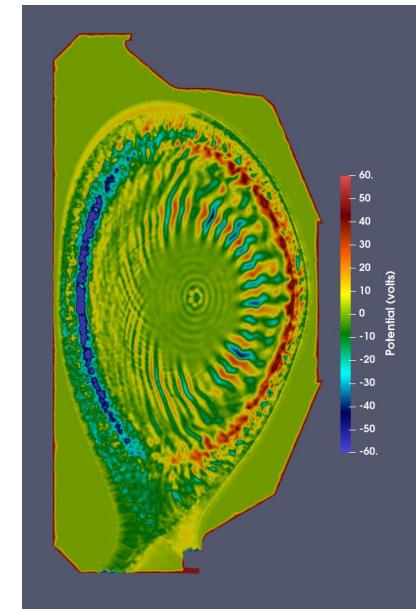
PUMIPic based XGCm Edge Plasma Code

XGCm is a version of XGC built on PUMIPic

- targeting execution of all operations on GPUs

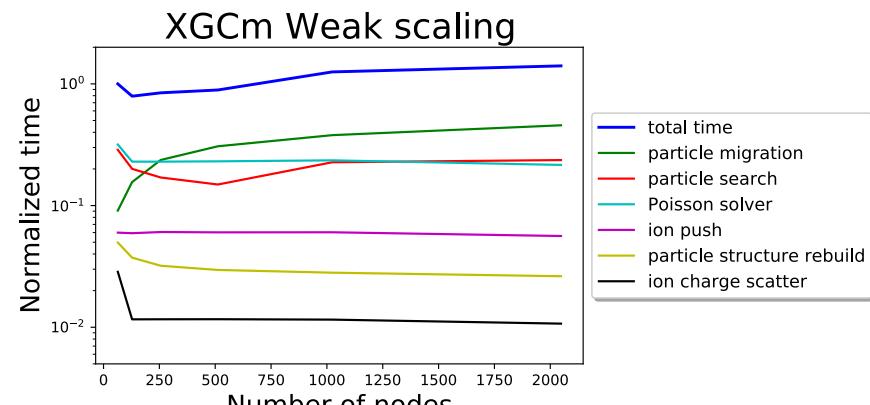
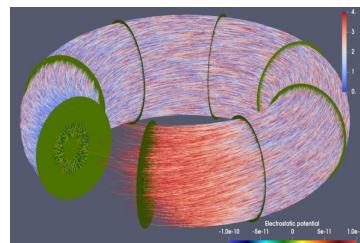
Testing of PUMIPic for use in XGC like push

- 2M elements, 1M vertices, 2 to 128 poloidal planes
- Pseudo push and particle-to-mesh gyro scatter
- Tested on up to 24,576 GPUs of Summit with 1.1 trillion particles, for 100 iterations: push, adjacency
- PUMIPic weak scaling up to 24576 GPUs (4096 nodes) with 48 million particles per GPU



XGCm status: All operations on GPU

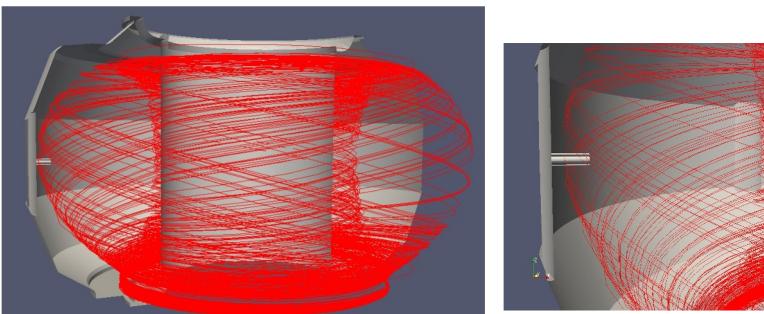
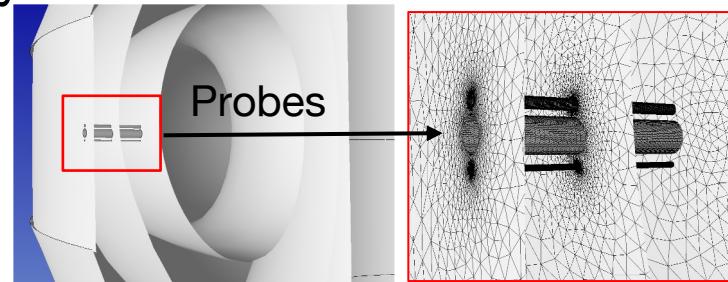
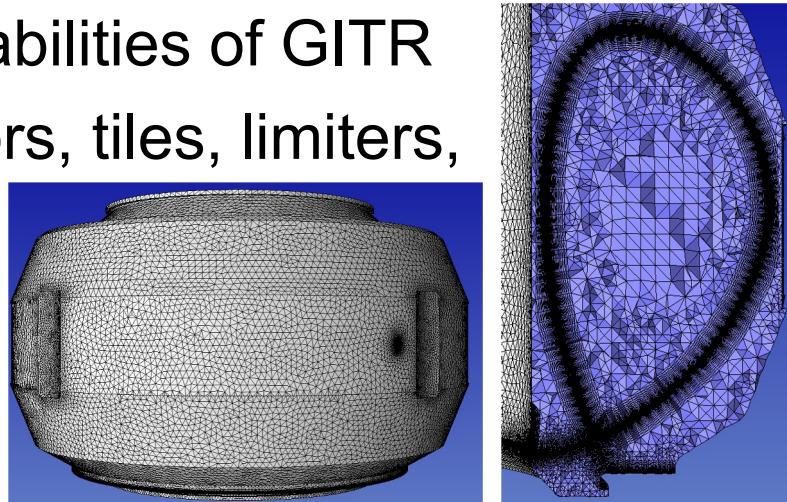
- Ion and electron scatter and push
- Electrostatic potential calculation
- Gyro-kinetic electric field calculation and gather
- Poisson solve



590,143 mesh elements, 20 million particles/GPU, $n_{planes} = n_{nodes}/8$

PUMIPic based GITRm Impurity Transport Code

- Incorporates impurity transport capabilities of GITR
- 3D mesh for cases including divertors, tiles, limiters, specific diagnostics/probes etc.
- Status
 - Physics equivalent to GITR
 - Particle initialization directly on 3D mesh
 - 3D mesh design/control including anisotropy to properly represent the background fields
 - Field transfer from SOLPS to 3D mesh
 - Non-uniform particle distribution
 - evolves quickly in time
 - Load balancing particles via EnGPar
 - Distance to boundary for sheath E field
 - Post-processing on 3D unstructured mesh



Run the latest Simmetrix and PUMI software on RPI systems

We will help you run the latest Simmetrix and PUMI model preparation, mesh generation, and adaptation tools on **your problem** using HPC systems at RPI.

Contact Cameron Smith in Slack, during Speed-Dating, or via email at smithc11@rpi.edu for more information.

